

MTH 602 Scientific Machine Learning

Homework 5

11/5/2025

S. M. Mahfuzul Hasan

02181922



I. PAPER & PENCIL WORK

1. Book problem 5.13:

The likelihood function can be written as,

$$p(\{\phi_n, \mathbf{t}_n\} | \{\pi_k\}) = \prod_{n=1}^N \prod_{k=1}^K \{p(\phi_n | \mathcal{C}_k) \pi_k\}^{t_{nk}} \quad (1)$$

Taking \ln in equation (1),

$$\begin{aligned} \ln [p(\{\phi_n, \mathbf{t}_n\} | \{\pi_k\})] &= \ln \left[\prod_{n=1}^N \prod_{k=1}^K \{p(\phi_n | \mathcal{C}_k) \pi_k\}^{t_{nk}} \right] \\ &= \sum_{n=1}^N \sum_{k=1}^K t_{nk} [\ln(p(\phi_n | \mathcal{C}_k)) + \ln \pi_k] \end{aligned} \quad (2)$$

To maximize equation (2) with respect to π_k , the constraint $\sum_{k=1}^K \pi_k = 1$ must be enforced. This can be done introducing a *Lagrange* multiplier. So, we need to maximize

$$\begin{aligned} \ln [p(\{\phi_n, \mathbf{t}_n\} | \{\pi_k\})] + \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) &= \sum_{n=1}^N \sum_{k=1}^K t_{nk} [\ln(p(\phi_n | \mathcal{C}_k)) + \ln \pi_k] \\ &\quad + \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) \end{aligned} \quad (3)$$

Now taking derivative of equation (3) and setting it to zero, we get

$$\begin{aligned} \frac{\partial}{\partial \pi_k} \left(\sum_{n=1}^N \sum_{k=1}^K t_{nk} [\ln(p(\phi_n | \mathcal{C}_k)) + \ln \pi_k] + \lambda \left(\sum_{k=1}^K \pi_k - 1 \right) \right) &= 0 \\ \Rightarrow 0 + \sum_{n=1}^N \frac{\partial}{\partial \pi_k} \left(t_{nk} \sum_{k=1}^K \ln \pi_k \right) + \lambda \frac{\partial}{\partial \pi_k} \left(\sum_{k=1}^K \pi_k \right) - 0 &= 0 \\ \Rightarrow \sum_{n=1}^N \frac{\partial}{\partial \pi_k} (t_{n1} \ln \pi_1 + \dots + t_{nk} \ln \pi_k + \dots + t_{nK} \ln \pi_K) + \lambda \frac{\partial}{\partial \pi_k} (\pi_1 + \dots + \pi_k + \dots + \pi_K) &= 0 \\ \Rightarrow \frac{\sum_{n=1}^N t_{nk}}{\pi_k} + \lambda = 0 &\quad [\text{since, } t_{nk} \text{ is a constant}] \\ \Rightarrow N_k = -\lambda \pi_k &\quad [\text{since, } \sum_{n=1}^N t_{nk} = N_k] \end{aligned} \quad (4)$$

Summing equation(4) on both sides with $\sum_{k=1}^K$,

$$\begin{aligned}\sum_{k=1}^K N_k &= -\lambda \sum_{k=1}^K \pi_k \\ \Rightarrow N &= -\lambda \cdot 1 \\ \Rightarrow \lambda &= -N\end{aligned}\tag{5}$$

Plugging λ value back in equation (4),

$$\begin{aligned}N_k &= -(-N)\pi_k \\ \Rightarrow \pi_k &= \frac{N_k}{N} \quad (\text{Showed})\end{aligned}\tag{6}$$

2. Book problem 5.14:

Given,

$$\begin{aligned}p(\phi_n | \mathcal{C}_k) &= \mathcal{N}(\phi_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}) \\ &= \frac{1}{2\pi^{D/2}|\boldsymbol{\Sigma}|^{1/2}} \exp \left\{ -\frac{1}{2}(\phi_n - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) \right\}\end{aligned}\tag{7}$$

Using equation (7) into equation (2),

$$\begin{aligned}&\ln [p(\{\phi_n, \mathbf{t}_n\} | \{\pi_k\})] \\ &= \sum_{n=1}^N \sum_{k=1}^K t_{nk} \left[\ln \left(\frac{1}{2\pi^{D/2}|\boldsymbol{\Sigma}|^{1/2}} \exp \left\{ -\frac{1}{2}(\phi_n - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) \right\} \right) + \ln \pi_k \right] \\ &= \sum_{n=1}^N \sum_{k=1}^K t_{nk} \left[-\ln(2\pi^{D/2}) - \frac{1}{2} \ln |\boldsymbol{\Sigma}| - \frac{1}{2}(\phi_n - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) + \ln \pi_k \right]\end{aligned}\tag{8}$$

Differentiating equation (8) with respect to $\boldsymbol{\mu}_k$ and setting it to zero, we get

$$\begin{aligned}\frac{\partial}{\partial \boldsymbol{\mu}_k} \left(\sum_{n=1}^N \sum_{k=1}^K t_{nk} \left[-\ln(2\pi^{D/2}) - \frac{1}{2} \ln |\boldsymbol{\Sigma}| - \frac{1}{2}(\phi_n - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) + \ln \pi_k \right] \right) &= 0 \\ \Rightarrow -0 - 0 - \sum_{n=1}^N t_{nk} \left(\frac{1}{2} \right) (-2) \boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) + 0 &= 0 \\ \Rightarrow \boldsymbol{\Sigma}^{-1} \sum_{n=1}^N t_{nk} (\phi_n - \boldsymbol{\mu}_k) &= 0 \\ \Rightarrow \sum_{n=1}^N t_{nk} \phi_n &= \boldsymbol{\mu}_k \sum_{n=1}^N t_{nk} \\ \Rightarrow \boldsymbol{\mu}_k &= \frac{1}{N_k} \sum_{n=1}^N t_{nk} \phi_n \quad [\text{since, } \sum_{n=1}^N t_{nk} = N_k] \\ &\quad (\text{Showed})\end{aligned}\tag{9}$$

Now, differentiating equation (8) with respect to $\boldsymbol{\Sigma}^{-1}$ and setting it to zero, we get

$$\begin{aligned}
& \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}} \left(\sum_{n=1}^N \sum_{k=1}^K t_{nk} \left[-\ln(2\pi^{D/2}) - \frac{1}{2} \ln |\boldsymbol{\Sigma}| - \frac{1}{2} (\phi_n - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) + \ln \pi_k \right] \right) = 0 \\
& \Rightarrow -0 - \frac{1}{2} \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}} \left(\sum_{n=1}^N \sum_{k=1}^K t_{nk} \ln |\boldsymbol{\Sigma}| \right) - \\
& \quad \frac{1}{2} \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}} \left(\sum_{n=1}^N \sum_{k=1}^K t_{nk} (\phi_n - \boldsymbol{\mu}_k)^T \boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) \right) + 0 = 0 \\
& \Rightarrow N(-\boldsymbol{\Sigma}^T) + \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}} \left(\sum_{n=1}^N \sum_{k=1}^K t_{nk} \text{Tr} \left(\boldsymbol{\Sigma}^{-1} (\phi_n - \boldsymbol{\mu}_k) (\phi_n - \boldsymbol{\mu}_k)^T \right) \right) = 0 \\
& \quad [\text{since, } \mathbf{x}^T A \mathbf{x} = \text{Tr}(A \mathbf{x} \mathbf{x}^T)] \\
& \Rightarrow -N\boldsymbol{\Sigma} + \sum_{n=1}^N \sum_{k=1}^K t_{nk} \left((\phi_n - \boldsymbol{\mu}_k) (\phi_n - \boldsymbol{\mu}_k)^T \right)^T = 0 \\
& \quad [\text{since, } \boldsymbol{\Sigma}^T = \boldsymbol{\Sigma}, \text{ and } \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}} (\text{Tr}(\boldsymbol{\Sigma}^{-1} A)) = A^T] \\
& \Rightarrow \boldsymbol{\Sigma} = \frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K t_{nk} (\phi_n - \boldsymbol{\mu}_k) (\phi_n - \boldsymbol{\mu}_k)^T \\
& = \sum_{k=1}^K \frac{N_k}{N} \cdot \frac{1}{N_k} \sum_{n=1}^N t_{nk} (\phi_n - \boldsymbol{\mu}_k) (\phi_n - \boldsymbol{\mu}_k)^T = \sum_{k=1}^K \frac{N_k}{N} \mathbf{S}_k \quad (\text{Showed})
\end{aligned}$$

II. USING GMM AS A GEOMETRIC AND PROBABILISTIC CLASSIFIER

A. Recap of the previous assignment (context)

Cont'd...

B. Setup and notation

Cont'd...

C. Data for this assignment

Ground truth mixture:

$$w = (0.2, 0.8); \quad \mu_1 = (0, 0), \quad \Sigma_1 = \begin{pmatrix} 1 & 0.8 \\ 0.8 & 1.5 \end{pmatrix}; \quad \mu_2 = (2, 2), \quad \Sigma_2 = \begin{pmatrix} 1.2 & -0.5 \\ -0.5 & 0.8 \end{pmatrix}$$

For sampling of the data, please refer to listing 1.




Figure 1: 3000 samples by first sampling the component index, and then sampling $\vec{x} \mid y = k \sim \mathcal{N}(\mu_k, \Sigma_k)$.

D. Training the Gaussian mixture model

1. Please refer to listing 1 and 2 for the fitted GMM.

Weights:

$$\hat{w} = \{0.83249319, 0.16750681\}$$

Means:

$$\hat{\mu} = \{(1.99485398, 1.95655549), (-0.07531723, -0.14512162)\}$$

Covariances:

$$\hat{\Sigma} = \left\{ \begin{pmatrix} 1.16217774 & -0.46118657 \\ -0.46118657 & 0.78502043 \end{pmatrix}, \begin{pmatrix} 0.86072986 & 0.61000376 \\ 0.61000376 & 1.1706272 \end{pmatrix} \right\}$$

2. Please refer to listing 1 and 2 for mapping.

Fitted GMM aligned to classes:

$$\hat{w} = (0.16750681, 0.83249319)$$

$$\hat{\mu}_1 = (-0.07531723, -0.14512162), \quad \hat{\Sigma}_1 = \begin{pmatrix} 0.86072986 & 0.61000376 \\ 0.61000376 & 1.1706272 \end{pmatrix}$$

$$\hat{\mu}_2 = (1.99485398, 1.95655549), \quad \hat{\Sigma}_2 = \begin{pmatrix} 1.16217774 & -0.46118657 \\ -0.46118657 & 0.78502043 \end{pmatrix}$$

3. Verification of the correctness of the GMM model has been done both quantitatively and qualitatively.





Figure 2: (a) Overlay of fitted samples on the true samples. (b) 1σ ellipse of Σ for fitted and true samples.

Firstly, The L_1 error norm for the weights and L_2 error norm for the means and covariances

were computed. The reported values are:

$$\begin{aligned}\|\hat{w}_1 - w_1\| &= 0.0325, & \|\hat{w}_2 - w_2\| &= 0.0325 \\ \|\hat{\mu}_1 - \mu_1\|_2 &= 0.1635, & \|\hat{\Sigma}_1 - \Sigma_1\|_2 &= 0.4468 \\ \|\hat{\mu}_2 - \mu_2\|_2 &= 0.0437, & \|\hat{\Sigma}_2 - \Sigma_2\|_2 &= 0.0669\end{aligned}$$

The fitted model almost captures the proportions of the mixture, showing 0.0325 absolute error for both the weights. The mean and covariance error for component 1 is higher comparatively than mean and covariance error for component 2. However, that is reasonable since the original weights were 0.2 for component 1, therefore it had very less sampled data points compared to component 2.

Then, the average log-likelihood per sample was calculated both for training and testing data. The values for training and testing data came out to be -3.0105 and -3.0615 , respectively. The higher log-likelihood on the training data indicates that the fitted GMM captures the training distribution well. The slightly lower test value is expected, as unseen data typically yield lower likelihoods. Since, the difference is only 0.05 ($\approx 1.7\%$), this indicates the good generalization behavior of the model, and that it is not prone to overfitting.

The visual check was done as a comparison between the fitted samples and true samples in figure 2. It shows that true and fitted samples span over the same regions. 1σ ellipses of Σ for fitted and true samples are almost exactly the same for component 2, while component 1 has a slight deviation towards right edge between the two ellipses, which can be attributed to the higher $L2$ norm of error Σ_1 . This is, again, due to small sample size of training data for component 1.

Lastly, the prediction accuracy of the model was checked on training and testing data using MAP decision rule. The accuracy shown on training and testing data are 95.12% and 95.33%, respectively which suggest that the model was able to predict seen and unseen data with the same level of accuracy. This is a sign of good generalization of the model. This indicates the capability of the model to separate classes convincingly, barring some misclassifications due to overlap between the tails of Gaussians, and limited sample size of component 1.

Confusion matrix for training data:

	Predicted: Class-1	Predicted: Class-2
Actual: Class-1	353	105
Actual: Class-2	12	1930

Confusion matrix for testing data:

	Predicted: Class-1	Predicted: Class-2
Actual: Class-1	82	26
Actual: Class-2	2	490

From the confusion matrices, it is clear that the majority of misclassification happens for component 1.

E. Geometric classification

- Decision boundary is reasonable for training data as can be seen from figure 3.




Figure 3: Visual check for training data and boundary.

- Decision boundary also seems reasonable for testing data as can be seen from figure 4.




Figure 4: Visual check for testing data, boundary, and misclassifications.

- Accuracy of the geometric classifier for testing data: 95.17%.

Confusion matrix:

	Predicted: Class-1	Predicted: Class-2
Actual: Class-1	86	22
Actual: Class-2	7	485

F. Probabilistic GMM classifier with thresholding & ROC

- Accuracy of the probabilistic classifier for testing data: 95.33%.

Confusion matrix:

	Predicted: Class-1	Predicted: Class-2
Actual: Class-1	82	26
Actual: Class-2	2	490

- Computed AUC for figure 5: 0.9622

Both the ROC curve and AUC indicate the very good separability of the classes for the probabilistic classifier. False positive rate increases a bit when True positive rate converges towards 1. However overall, the classifier shows really good capability in identifying the classes.




Figure 5: ROC curve.

G. Comparison: Probabilistic vs geometric classifier

- Both classifiers show similar accuracy, but probabilistic classifier shows 0.16% better accuracy compared to geometric classifier. From the confusion matrices of the respective classifiers, and taking the MAP classification from GMM into account, geometric classifier misclassifies class 2 more than probabilistic classifier, where probabilistic classifier is fully in agreement with MAP classification. For a large sample size as class 2, there should be least number of misclassifications for class 2. Geometric classifier shows comparatively bad performance

in that case. However, geometric classifier shows better performance for classifying class 1, but since probabilistic classifier relies on posterior probabilities rather than just a fixed perpendicular bisector, it has smoother transition between classes, and higher accuracy overall. Hence, it is slightly better for overlapping regions.

- (ii) The dataset constitute a comparatively easy classification task because of highly separable classes due to limited overlap, and thus higher accuracy of the classifiers. High AUC and very good generalization also point to that.
- (iii) Class 1 is more difficult to predict for its small sample size due to small weight. This can also be inferred from the confusion matrices, where the misclassifications largely happened for class 1. Class 1 has larger variance, and spreads well into class 2 region, making it difficult for the classifiers to detect them accurately.

```

1  from sklearn.mixture import GaussianMixture
2  from sklearn.metrics import confusion_matrix, accuracy_score, roc_curve, auc
3  from sklearn.model_selection import train_test_split
4  import numpy as np
5  import matplotlib.pyplot as plt
6  import matplotlib as mpl
7  from matplotlib.patches import Ellipse
8
9  # ===== B. Setup and notation =====
10 def sample_2d_gaussian(mu, Sigma, n, rng=None):
11     if rng is None:
12         rng = np.random.default_rng()
13     return rng.multivariate_normal(mean=mu, cov=Sigma, size=n)
14
15
16 # ===== C. Data for this assignment =====
17 # setting parameters
18 rng = np.random.default_rng(29)
19 n = 3000
20
21 # ground-truth
22 w_true = np.array([0.2, 0.8])
23 mu1 = np.array([0.0, 0.0])
24 Sigma1 = np.array([[1.0, 0.8], [0.8, 1.5]])
25 mu2 = np.array([2.0, 2.0])
26 Sigma2 = np.array([[1.2, -0.5], [-0.5, 0.8]])
27
28 # initialization
29 y = rng.choice([1, 2], size=n, p=[0.2, 0.8])
30 x_ = np.zeros((n, 2))
31
32 # sampling
33 n1 = np.sum(y == 1); n2 = np.sum(y == 2)
34 x_[y == 1] = sample_2d_gaussian(mu1, Sigma1, n1, rng)
35 x_[y == 2] = sample_2d_gaussian(mu2, Sigma2, n2, rng)
36
37 # parameters for plotting
38 plt.rcParams["font.family"] = "serif"
39 plt.rcParams["font.serif"] = ["CMU Serif"]
40 plt.rcParams["mathtext.fontset"] = "cm"
41 plt.rcParams["font.size"] = 20
42 mpl.rcParams["axes.unicode_minus"] = False
43
44 # samples plot

```

```

45 fig, ax = plt.subplots(figsize=(15, 6))
46
47 ax.scatter(x_[y==1, 0], x_[y==1, 1], color="blue", label="Component 1")
48 ax.scatter(x_[y==2, 0], x_[y==2, 1], color="red", label="Component 2")
49
50 plt.xlabel(r'$x_1$')
51 plt.ylabel(r'$x_2$')
52 plt.xlim(-5.5,6)
53 plt.ylim(-3, 4.5)
54 plt.legend(loc="upper left", frameon=False)
55 plt.tick_params(axis="both", which="both", direction="in")
56 plt.savefig(f'x_sample.pdf', dpi=1080)
57 plt.show()
58
59 # 80-20 split
60 x_train, x_test, y_train, y_test = train_test_split(x_, y, test_size=0.2,
61 random_state=29)
62
63 # ===== D. Training the Gaussian mixture model =====
64 # 1. GMM fit
65 g2 = GaussianMixture(n_components=2, covariance_type='full',
66 reg_covar=1e-6, n_init=10, random_state=29)
66 g2.fit(x_train)
67
68 print("\n1. GMM fit:\n\nFor K=2, ")
69 print("\nWeights:", g2.weights_)
70 print("\nMeans:", g2.means_)
71 print("\nCovariances:", g2.covariances_)
72
73 # 2. Mapping components to true classes
74 M = g2.means_
75 d0 = np.linalg.norm(M[0] - mu1) + np.linalg.norm(M[1] - mu2)
76 d1 = np.linalg.norm(M[1] - mu1) + np.linalg.norm(M[0] - mu2)
77 if d0 <= d1:
78     class_to_component = {1: 0, 2: 1}
79     component_to_class = {0: 1, 1: 2}
80 else:
81     class_to_component = {1: 1, 2: 0}
82     component_to_class = {1: 1, 0: 2}
83
84 print("\n2. Mapping components to classes:\n\n", class_to_component)
85 print("component_to_class:", component_to_class)
86
87 # reordering components to match classes
88 w_hat1 = g2.weights_[class_to_component[1]]
89 w_hat2 = g2.weights_[class_to_component[2]]
90 mu_hat1 = g2.means_[class_to_component[1]]
91 mu_hat2 = g2.means_[class_to_component[2]]
92 Sigma_hat1 = g2.covariances_[class_to_component[1]]
93 Sigma_hat2 = g2.covariances_[class_to_component[2]]
94
95 print("\nTrue and aligned fitted GMM:")
96 print("\nTrue weights:", w_true)
97 print("Aligned fitted weights:", [w_hat1, w_hat2])
98 print("\nTrue means:\n", np.vstack([mu1, mu2]))
99 print("\nAligned fitted means:\n", np.vstack([mu_hat1, mu_hat2]))
100 print("\nTrue covariances:")
101 print(Sigma1, "\n\n", Sigma2)

```

```

102 print("\nAligned fitted covariances:")
103 print(Sigma_hat1, "\n\n", Sigma_hat2)
104
105 # 3. Verification of the correctness of the GMM model
106 # L2 norm of errors
107 w1_err = abs(w_hat1 - w_true[0])
108 w2_err = abs(w_hat2 - w_true[1])
109 mu_err1 = np.linalg.norm(mu_hat1 - mu1, ord=2)
110 mu_err2 = np.linalg.norm(mu_hat2 - mu2, ord=2)
111 Sig_err1 = np.linalg.norm(Sigma_hat1 - Sigma1, ord=2)
112 Sig_err2 = np.linalg.norm(Sigma_hat2 - Sigma2, ord=2)
113
114 print(f"\n3. Verification of GMM model:")
115 print(f"||w_1(fit) - w_1(true)|| = {w1_err:.4f}")
116 print(f"||w_2(fit) - w_2(true)|| = {w2_err:.4f}")
117 print(f"||mu_1(fit) - mu_1(true)||_2 = {mu_err1:.4f}")
118 print(f"||mu_2(fit) - mu_2(true)||_2 = {mu_err2:.4f}")
119 print(f"||Sigma_1(fit) - Sigma_1(true)||_2 = {Sig_err1:.4f}")
120 print(f"||Sigma_2(fit) - Sigma_2(true)||_2 = {Sig_err2:.4f}")
121
122 # average log-likelihood on training and test data
123 train_logl = g2.score(x_train)
124 test_logl = g2.score(x_test)
125
126 print(f"\nAverage log-likelihood per sample (train): {train_logl:.4f}")
127 print(f"Average log-likelihood per sample (test): {test_logl:.4f}")
128
129 # Draw samples from the fitted GMM (same number as true data)
130 x_fit, y_fit = g2.sample(n)
131
132 # true and fitted GMMs plot
133 fig, ax = plt.subplots(figsize=(15, 6))
134
135 ax.scatter(x_train[y_train == 1, 0], x_train[y_train == 1, 1],
136             color="blue", alpha=0.6, label="True (component 1)")
137 ax.scatter(x_train[y_train == 2, 0], x_train[y_train == 2, 1],
138             color="red", alpha=0.6, label="True (component 2)")
139
140 ax.scatter(x_fit[y_fit == 0, 0], x_fit[y_fit == 0, 1], color="orange", alpha=0.7,
141             marker='x', label="Fitted (component 1)")
142 ax.scatter(x_fit[y_fit == 1, 0], x_fit[y_fit == 1, 1], color="green", alpha=0.7,
143             marker='x', label="Fitted (component 2)")
144
145 plt.xlabel(r"$x_1$")
146 plt.ylabel(r"$x_2$")
147 plt.xlim(-5.5, 6)
148 plt.ylim(-3, 4.5)
149 plt.legend(loc="upper left", frameon=False)
150 plt.tick_params(axis="both", which="both", direction="in")
151 plt.savefig(f'x_sample_fit.pdf', dpi=1080)
152 plt.show()
153
154 # 1$ \sigma$ ellipse
155 fig, ax = plt.subplots(figsize=(15, 6))
156
157 # true ellipses
158 for k, (mu_true, Sigma_true, color) in enumerate(
159     [(mu1, Sigma1, "blue"), (mu2, Sigma2, "red")], start=1):

```

```

161 # compute eigenvalues & eigenvectors
162 eigen_val, eigen_vec = np.linalg.eigh(Sigma_true)
163
164 # sort in descending order
165 order = np.argsort(eigen_val)[::-1]
166 eigen_val = eigen_val[order]
167 eigen_vec = eigen_vec[:, order]
168
169 # rotation angle
170 theta = np.degrees(np.arctan2(*eigen_vec[:, 0][::-1])) % 180
171
172 # major and minor axes
173 width, height = 2 * np.sqrt(eigen_val)
174
175 # print for reference (optional)
176 print(f"\nTrue Component {k}:")
177 print(f"  Eigenvalues = {eigen_val}")
178 print(f"  Rotation = {theta:.2f} degree")
179
180 ell = Ellipse(xy=mu_true, width=width, height=height, angle=theta, ls="--",
181                 edgecolor=color, facecolor='none', lw=3,
182                 label=f"True ellipse (component {k})")
183 ax.add_patch(ell)
184
185 # fitted ellipses
186 for k, (mu_hat, Sigma_hat, color) in enumerate(
187     [(mu_hat1, Sigma_hat1, "green"), (mu_hat2, Sigma_hat2, "orange")], start
188     =1):
189
190     # compute eigenvalues & eigenvectors
191     eigen_val, eigen_vec = np.linalg.eigh(Sigma_hat)
192
193     # sort in descending order
194     order = np.argsort(eigen_val)[::-1]
195     eigen_val = eigen_val[order]
196     eigen_vec = eigen_vec[:, order]
197
198     # rotation angle
199     theta = np.degrees(np.arctan2(*eigen_vec[:, 0][::-1])) % 180
200
201     # major and minor axes
202     width, height = 2 * np.sqrt(eigen_val)
203
204     # print for reference (optional)
205     print(f"\nFitted Component {k}:")
206     print(f"  Eigenvalues = {eigen_val}")
207     print(f"  Rotation = {theta:.2f} degree")
208
209     ell = Ellipse(xy=mu_hat, width=width, height=height, angle=theta, ls="--",
210                   edgecolor=color, facecolor='none', lw=4,
211                   label=f"Fitted ellipse (component {k})")
212     ax.add_patch(ell)
213
214 plt.xlabel(r"$x_1$")
215 plt.ylabel(r"$x_2$")
216 plt.xlim(-5.5, 6)
217 plt.ylim(-3, 4.5)
218 plt.legend(loc="upper left", frameon=False)
219 plt.tick_params(axis="both", which="both", direction="in")

```

```

219 plt.savefig(f'ellipse.pdf', dpi=1080)
220 plt.show()
221
222 # training and testing data prediction accuracy
223 comp_pred_train = g2.predict(x_train)
224 y_pred_train = np.array([component_to_class[c] for c in comp_pred_train])
225 train_acc = np.mean(y_pred_train == y_train)
226
227 comp_pred_test = g2.predict(x_test)
228 y_pred_test = np.array([component_to_class[c] for c in comp_pred_test])
229 test_acc = np.mean(y_pred_test == y_test)
230
231 print(f"\nTraining accuracy: {train_acc*100:.2f}%")
232 print(f"Test accuracy: {test_acc*100:.2f}%")
233
234 print("\nConfusion matrix (train):")
235 print(confusion_matrix(y_train, y_pred_train, labels=[1, 2]))
236
237 print("\nConfusion matrix (test):")
238 print(confusion_matrix(y_test, y_pred_test, labels=[1, 2]))
239
240
241 # ===== E. Geometric classification =====
242 # 1. Visual check for training data and boundary
243 # geometric parameters
244 r_hat = (mu_hat2 - mu_hat1) / np.linalg.norm(mu_hat2 - mu_hat1)
245 m = 0.5 * (mu_hat1 + mu_hat2)
246
247 def classify_geom(x):
248     return np.where(np.dot(x - m, r_hat) < 0, 1, 2)
249
250 # decision boundary
251 x_vals = np.linspace(np.min(x_train[:,0]) - 1, np.max(x_train[:,0]) + 1, 200)
252 y_vals = m[1] - (r_hat[0]/r_hat[1]) * (x_vals - m[0])
253
254 # plot training boundary and data
255 fig, ax = plt.subplots(figsize=(15, 6))
256 ax.scatter(x_train[y_train==1,0], x_train[y_train==1,1], c='blue',
257             alpha=0.5, label='Class 1 (train)')
258 ax.scatter(x_train[y_train==2,0], x_train[y_train==2,1], c='red',
259             alpha=0.5, label='Class 2 (train)')
260 ax.plot(x_vals, y_vals, 'k--', lw=3, label='Decision boundary')
261
262 plt.xlabel(r'$x_1$');
263 plt.ylabel(r'$x_2$')
264 plt.legend(loc="upper right", frameon=False)
265 plt.tick_params(axis="both", which="both", direction="in")
266 plt.savefig(f'train_bndry.pdf', dpi=1080)
267 plt.show()
268
269 # 2. Visual check for testing data, boundary, and misclassification
270 # classify testing data
271 y_geom_test = classify_geom(x_test)
272
273 # plot testing decision boundary and data
274 fig, ax = plt.subplots(figsize=(15, 6))
275 ax.scatter(x_test[y_test==1,0], x_test[y_test==1,1], c='blue', alpha=0.5,
276             label='Class 1 (test)')
277 ax.scatter(x_test[y_test==2,0], x_test[y_test==2,1], c='red', alpha=0.5,

```

```

278         label='Class 2 (test)')
279 ax.plot(x_vals, y_vals, 'k--', lw=2, label='Decision boundary')
280
281 mis_idx = (y_geom_test != y_test)
282 ax.scatter(x_test[mis_idx,0], x_test[mis_idx,1], facecolors='none',
283             edgecolors='black', lw=1.5, s=80, label='Misclassified')
284
285 plt.xlabel(r'$x_1$');
286 plt.ylabel(r'$x_2$')
287 plt.legend(loc="lower left", frameon=False)
288 plt.tick_params(axis="both", which="both", direction="in")
289 plt.savefig(f'test_bndry.pdf', dpi=1080)
290 plt.show()
291
292 # 3. Confusion matrix (test)
293 confuse_geom = confusion_matrix(y_test, y_geom_test, labels=[1, 2])
294 accuracy_geom = accuracy_score(y_test, y_geom_test)
295 print("\n3. Accuracy and confusion matrix of geometric classifier:\n\nAccuracy (test): {:.2f}%".format(accuracy_geom*100))
296 print("Confusion matrix (test):")
297 print(confuse_geom)
298
299 # ===== F. Probabilistic classifier with thresholding & ROC =====
300 # Compute posterior probabilities for each class using the fitted GMM
301 proba = g2.predict_proba(x_test)
302
303 # identify the fitted GMM component corresponding to true class 2
304 comp_idx_for_class2 = class_to_component[2]
305
306 # extract posterior probability
307 proba_class2 = proba[:, comp_idx_for_class2]
308
309 # classification based on threshold
310 tau = 0.5
311 y_prob_test = np.where(proba_class2 >= tau, 2, 1)
312
313 # 1. Confusion matrix (test)
314 accuracy_prob = accuracy_score(y_test, y_prob_test)
315 confuse_prob = confusion_matrix(y_test, y_prob_test, labels=[1, 2])
316
317 print(f"\n1. Accuracy and confusion matrix of probabilistic classifier:\n\nAccuracy (test): {accuracy_prob*100:.2f}%")
318 print("Confusion matrix (test):")
319 print(confuse_prob)
320
321 # 2. ROC curve and AUC
322 # convert true labels {1,2} to {0,1}
323 y_test_b = (y_test == 2).astype(int)
324
325 # compute ROC curve points across thresholds
326 f_p, t_p, thresholds = roc_curve(y_test_b, proba_class2)
327
328 # compute AUC
329 roc_auc = auc(f_p, t_p)
330
331 print(f"\n2. ROC curve and AUC:\n\nAUC= {roc_auc:.4f}")
332
333 # plot ROC curve
334 fig, ax = plt.subplots(figsize=(8, 8))

```

```

335 ax.plot(f_p, t_p, color='darkorange', lw=3, label=f'ROC curve (AUC = {roc_auc:.4f})')
336 ax.plot([0, 1], [0, 1], color='blue', lw=2, ls='--', label="Chance")
337
338 plt.xlabel('False Positive Rate (FPR)')
339 plt.ylabel('True Positive Rate (TPR)')
340 plt.legend(frameon=False)
341 plt.tick_params(axis="both", which="both", direction="in")
342 plt.savefig(f'roc.pdf', dpi=1080)
343 plt.show()

```

Listing 1: *gmm.py*

```

1 1. GMM fit:
2
3 For K=2,
4
5 Weights: [0.83249319 0.16750681]
6
7 Means: [[ 1.99485398  1.95655549]
8 [-0.07531723 -0.14512162]]
9
10 Covariances: [[[ 1.16217774 -0.46118657]
11 [-0.46118657  0.78502043]]
12
13 [[ 0.86072986  0.61000376]
14 [ 0.61000376  1.1706272 ]]]
15
16 2. Mapping components to classes:
17
18 class_to_component: {1: 1, 2: 0}
19 component_to_class: {1: 1, 0: 2}
20
21 True and aligned fitted GMM:
22
23 True weights: [0.2 0.8]
24 Aligned fitted weights: [0.1675068064372643, 0.8324931935627358]
25
26 True means:
27 [[0. 0.]
28 [2. 2.]]
29
30 Aligned fitted means:
31 [[-0.07531723 -0.14512162]
32 [ 1.99485398  1.95655549]]
33
34 True covariances:
35 [[1. 0.8]
36 [0.8 1.5]]
37
38 [[ 1.2 -0.5]
39 [-0.5  0.8]]
40
41 Aligned fitted covariances:
42 [[0.86072986 0.61000376]
43 [0.61000376 1.1706272 ]]
44
45 [[ 1.16217774 -0.46118657]
46 [-0.46118657  0.78502043]]

```

```

47
48 3. Verification of GMM model:
49
50 ||w_1(fit) - w_1(true)|| = 0.0325
51 ||w_2(fit) - w_2(true)|| = 0.0325
52 ||mu_1(fit) - mu_1(true)||_2 = 0.1635
53 ||mu_2(fit) - mu_2(true)||_2 = 0.0437
54 ||Sigma_1(fit) - Sigma_1(true)||_2 = 0.4468
55 ||Sigma_2(fit) - Sigma_2(true)||_2 = 0.0669
56
57 Average log-likelihood per sample (train): -3.0105
58 Average log-likelihood per sample (test): -3.0615
59
60 True Component 1:
61   Eigenvalues = [2.08815273 0.41184727]
62   Rotation = 53.68 degree
63
64 True Component 2:
65   Eigenvalues = [1.53851648 0.46148352]
66   Rotation = 145.90 degree
67
68 Fitted Component 1:
69   Eigenvalues = [1.64505416 0.38630291]
70   Rotation = 52.13 degree
71
72 Fitted Component 2:
73   Eigenvalues = [1.47185099 0.47534718]
74   Rotation = 146.12 degree
75
76 Training accuracy: 95.12%
77 Test accuracy: 95.33%
78
79 Confusion matrix (train):
80 [[ 353  105]
81  [ 12 1930]]
82
83 Confusion matrix (test):
84 [[ 82  26]
85  [  2 490]]
86
87 3. Accuracy and confusion matrix of geometric classifier:
88
89 Accuracy (test): 95.17%
90 Confusion matrix (test):
91 [[ 86  22]
92  [  7 485]]
93
94 1. Accuracy and confusion matrix of probabilistic classifier:
95
96 Accuracy (test): 95.33%
97 Confusion matrix (test):
98 [[ 82  26]
99  [  2 490]]
100
101 2. ROC curve and AUC:
102
103 AUC= 0.9622

```

Listing 2: Output terminal for *gmm.py*

III. PARTICLE PHYSICS: CLASSIFICATION OF A PARENT PARTICLE FROM TWO-BODY DECAY KINEMATICS

Four physics-informed features were taken into account for the solution of the problem.

Invariant mass:

For $c = 1$,

$$m = \sqrt{(E_1 + E_2)^2 - \left((P_{x1} + P_{x2})^2 + (P_{y1} + P_{y2})^2 + (P_{z1} + P_{z2})^2 \right)}$$

Opening angle:

$$\cos\theta = \frac{P_{x1} \cdot P_{x2} + P_{y1} \cdot P_{y2} + P_{z1} \cdot P_{z2}}{\sqrt{\left(P_{x1}^2 + P_{y1}^2 + P_{z1}^2\right) \left(P_{x2}^2 + P_{y2}^2 + P_{z2}^2\right)}}$$

Transverse momentum:

$$P_T = \sqrt{(P_{x1} + P_{x2})^2 + (P_{y1} + P_{y2})^2}$$

Energy asymmetry:

$$E_{asym} = \frac{E_1 - E_2}{E_1 + E_2}$$

where, subscripts 1 and 2 refer to daughter 1 and 2, respectively.




Figure 6: ROC curve for the classifier.

After computing these physical features, the training data were split 80-20 for training and validation. The features were scaled to mean 0 and variance 1. Then, the probabilistic classifier *Logistic*

Regression model was used, where

$$p(B|\mathbf{x}) = \frac{1}{1 + \exp(-\phi(\mathbf{x}; \boldsymbol{\omega}))}$$

The accuracy of the model was computed for validation set: 99.60%.

And, the confusion matrix:

	Predicted: Particle A	Predicted: Particle B
Actual: Particle A	1200	5
Actual: Particle B	11	2784

The prediction performance of the classifier is pretty good. That is further be verified from the ROC curve in figure 6. The AUC is 0.99925237, which is very close to 1.

For the test set, predicted count of parent types:

$$N_A_{pred} = 155 \quad N_B_{pred} = 345$$

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_score,
    roc_curve
7 import os
8
9 # very small constant to avoid division by zero
10 epsi = 1e-15
11
12 # load training data
13 data = np.genfromtxt("train_hw5.csv", delimiter=",", skip_header=1, dtype=str)
14
15 # separate features and labels
16 x_raw = data[:, :-1].astype(float)      # first 8 columns
17 labels = data[:, -1]                    # last column
18
19 # map A to 0, B to 1
20 y = np.where(labels == "A", 0, 1)
21
22 # unpack columns
23 E1, px1, py1, pz1, E2, px2, py2, pz2 = [x_raw[:, i] for i in range(8)]
24
25 # physics-informed features
26 E = E1 + E2 # total energy of the daughters
27 px = px1 + px2 # total x-momentum
28 py = py1 + py2 # total y-momentum
29 pz = pz1 + pz2 # total z-momentum
30
31 # total invariant mass
32 m2 = E**2 - (px**2 + py**2 + pz**2)
33 m = np.sqrt(np.clip(m2, 0, None))
34

```

```

35 # opening angle
36 p1 = np.sqrt(px1**2 + py1**2 + pz1**2)
37 p2 = np.sqrt(px2**2 + py2**2 + pz2**2)
38 cos_open = (px1*px2 + py1*py2 + pz1*pz2) / np.clip(p1*p2, epsi, None)
39
40 # transverse momentum
41 p_t = np.sqrt(px**2 + py**2)
42
43 # energy asymmetry
44 E_asym = (E1 - E2) / np.clip(E1 + E2, epsi, None)
45
46 # feature matrix
47 x = np.column_stack([m, p_t, cos_open, E_asym])
48
49 # 80-20 split and scaling
50 x_train, x_valid, y_train, y_valid = train_test_split(x, y, test_size=0.2,
51                                         random_state=29, stratify=y)
52 scaler = StandardScaler()
53 x_train_s = scaler.fit_transform(x_train)
54 x_valid_s = scaler.transform(x_valid)
55
56 # logistic regression
57 clf = LogisticRegression(max_iter=1000, C=2.0)
58 clf.fit(x_train_s, y_train)
59
60 # evaluate validation data
61 y_pred = clf.predict(x_valid_s)
62 y_prob = clf.predict_proba(x_valid_s)[:, 1]
63 acc = accuracy_score(y_valid, y_pred)
64 cm = confusion_matrix(y_valid, y_pred)
65 auc = roc_auc_score(y_valid, y_prob)
66
67 print(f"\nAccuracy: {acc*100:.2f}%")
68 print("\nConfusion matrix:\n", cm)
69 print(f"\nAUC: {auc:.8f}")
70
71 f_p, t_p, _ = roc_curve(y_valid, y_prob)
72
73 # ROC plot
74 fig, ax = plt.subplots(figsize=(8, 8))
75 ax.plot(f_p, t_p, color='darkorange', lw=3, label=f'ROC curve (AUC = {auc:.8f})')
76 ax.plot([0, 1], [0, 1], color='blue', lw=2, ls='--', label="Chance")
77
78 plt.xlabel('False Positive Rate (FPR)')
79 plt.ylabel('True Positive Rate (TPR)')
80 plt.legend(frameon=False)
81 plt.tick_params(axis="both", which="both", direction="in")
82 plt.savefig(f'roc_particle.pdf', dpi=1080)
83 plt.show()
84
85 # predict test data
86 test_data = np.genfromtxt("test_hw5.csv", delimiter=",", skip_header=1)
87 E1, px1, py1, pz1, E2, px2, py2, pz2 = [test_data[:, i] for i in range(8)]
88
89 E = E1 + E2; px = px1 + px2; py = py1 + py2; pz = pz1 + pz2
90 m = np.sqrt(np.clip(E**2 - (px**2 + py**2 + pz**2), 0, None))
91 p_t = np.sqrt(px**2 + py**2)
92 p1 = np.sqrt(px1**2 + py1**2 + pz1**2)
93 p2 = np.sqrt(px2**2 + py2**2 + pz2**2)

```

```

94 cos_open = (px1*px2 + py1*py2 + pz1*pz2) / np.clip(p1*p2, epsi, None)
95 E_asym = (E1 - E2) / np.clip(E1 + E2, epsi, None)
96
97 x_test = np.column_stack([m, p_t, cos_open, E_asym])
98 x_test_s = scaler.transform(x_test)
99
100 y_test_pred = clf.predict(x_test_s)
101
102
103 # probabilistic outputs from the logistic regression
104 prob = clf.predict_proba(x_test_s)[:, 1]
105 N_test = len(prob)
106
107 # soft counts
108 N_B_pred_soft = np.sum(prob)
109 N_A_pred_soft = N_test - N_B_pred_soft
110
111 # hard counts
112 N_B_pred_hard = np.sum(y_test_pred == 1)
113 N_A_pred_hard = np.sum(y_test_pred == 0)
114
115 print("\nPredicted parent counts:")
116 print(f"\nSoft count:")
117 print(f"\nN_A_pred = {N_A_pred_soft:.2f}, N_B_pred = {N_B_pred_soft:.2f}")
118 print(f"\nHard count (threshold 0.5):")
119 print(f"\nN_A_pred = {N_A_pred_hard}, N_B_pred = {N_B_pred_hard}")
120 print(f"\nTotal parents in test set: N_test = {N_test}")
121
122 y_cd = np.where(y_test_pred == 1, 2, 1).astype(int)
123
124 # write labels to text file
125 if os.path.exists("Mahfuzul_hw5.txt"):
126     os.remove("Mahfuzul_hw5.txt")
127
128 with open("Mahfuzul_hw5.txt", "w", encoding="utf-8") as f:
129     for val in y_cd:
130         f.write(f"{val}\n")
131
132 print(f"\nSaved Mahfuzul_hw5.txt with {len(y_cd)} predictions.")
133
134 # check file
135 y_check = np.loadtxt("Mahfuzul_hw5.txt", dtype=int)
136 print(y_check[:10])

```

Listing 3: *particle.py*

```
1 Accuracy: 99.60%
2
3 Confusion matrix:
4 [[1200      5]
5 [   11 2784]]
6
7 AUC: 0.99925237
8
9 Predicted parent counts:
10
11 Soft count:
12
13 N_A_pred = 155.21, N_B_pred = 344.79
14
15 Hard count (threshold 0.5):
16
17 N_A_pred = 155, N_B_pred = 345
18
19 Total parents in test set: N_test = 500
20
21 Saved Mahfuzul_hw5.txt with 500 predictions.
22 [2 2 1 2 2 1 2 2 2 2]
```

Listing 4: Output terminal for *particle.py*