

MTH 602 Scientific Machine Learning

Homework 7

12/9/2025

S. M. Mahfuzul Hasan

02181922



I. CNNS FOR SIGNAL DETECTION: CAN A NONLINEAR CLASSIFIER BEAT THE MATCHED FILTER?

B. Matched-filter recap and extended study

- For each amplitude $A \in \{0, 3\}$, 20000 independent realizations of \mathbf{y} are generated and corresponding ρ values are computed.

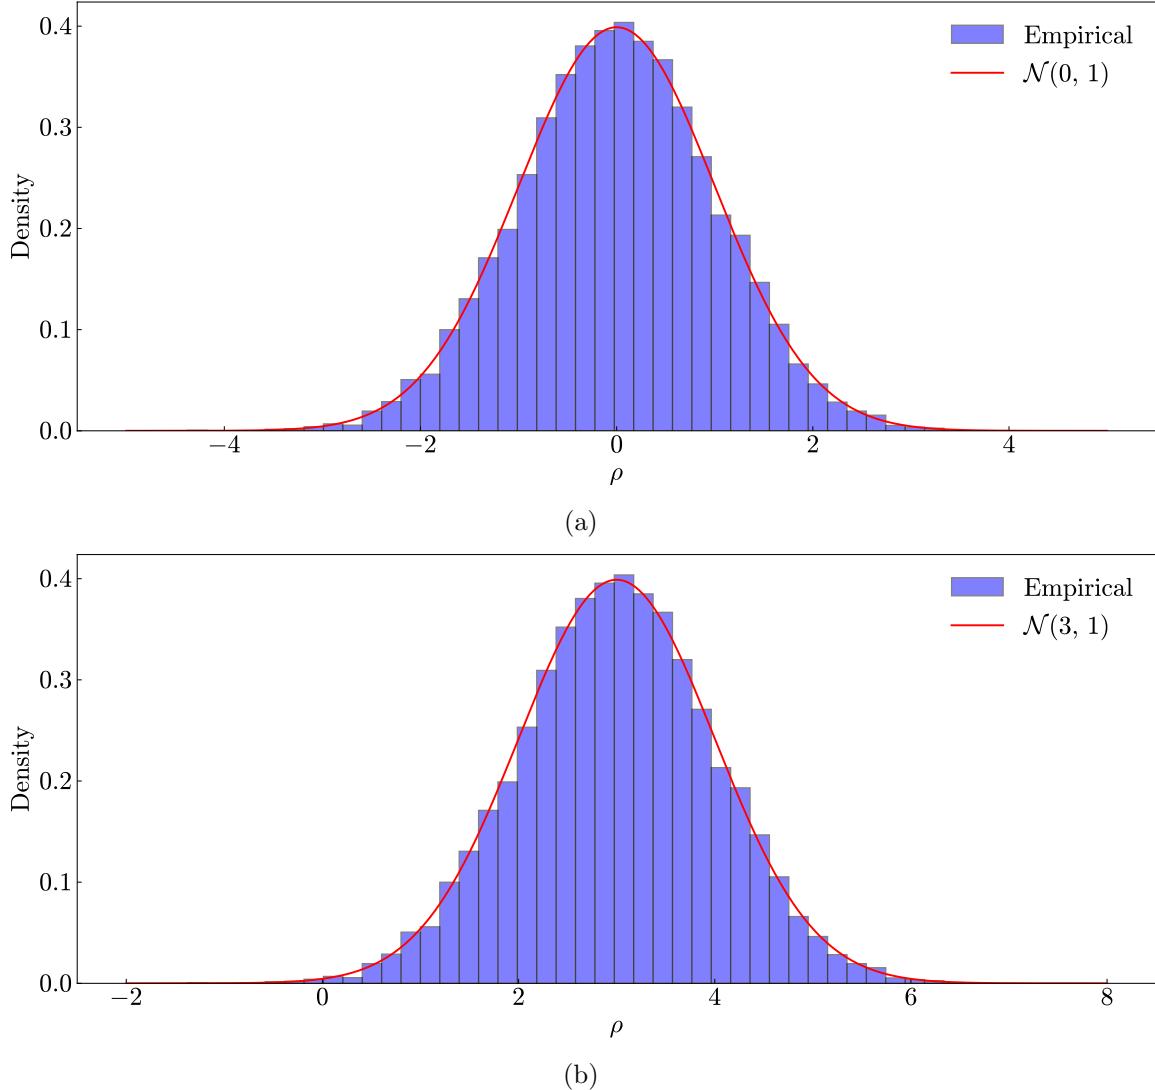


Figure 1: Empirical histogram of ρ and theoretical Gaussian PDF for (a) $A = 0$ and (b) $A = 3$.

The empirical distributions of ρ and theoretical Gaussian PDFs in fig. 1 for both the amplitude cases show very good agreement.

- Please refer to listing 1 for `filter_classifier`.

To verify that the function is working properly, first, it is checked whether there is any mismatch between the prediction of class between theory and `filter_classifier`.

The number of mismatched prediction = 0

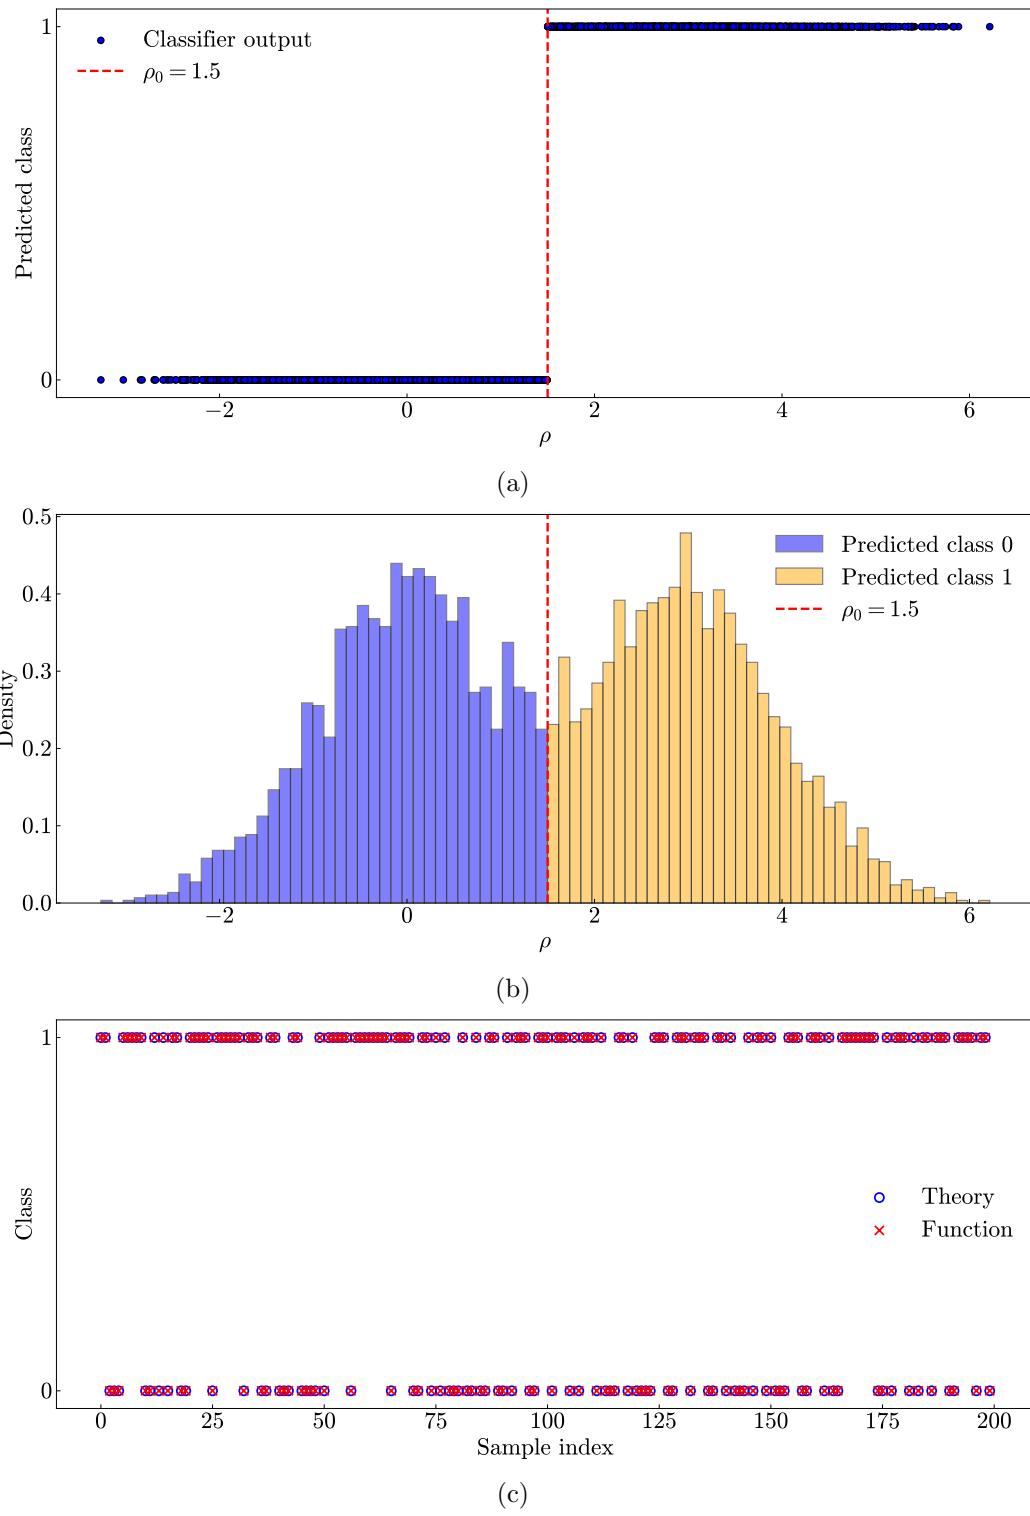


Figure 2: (a) Scatter plot of predicted class against ρ , (b) PDF of predicted class based on ρ_0 , and (c) comparison between theory and `filter_classifier` for the prediction of class for the first 200 samples.

Then, the predicted class and their PDF are plotted against the detection statistics in fig. 2a and 2b respectively. Both show clear separation of "signal" and "no-signal" based on the threshold value ρ_0 . Along with that, the comparison between the theoretical estimation of classes and estimation of classes by filter_classifier shows complete agreement in fig. 2c.

C. Building a labeled dataset for a CNN

1. Please refer to listing 1 for the construction of a labeled dataset of time-series

$$\left\{ \left(\mathbf{y}^{(k)}, c^{(k)} \right) \right\}_{k=1}^M, \quad c^{(k)} \in \{0, 1\}$$

The number of samples is set at 20000 (found to be sufficiently optimal), and the splitting renders

Number of no-signal examples: 9933 (49.665%)

Number of signal examples: 10067 (50.335%)

which shows almost 50% class balance.

2. The dataset is split into 80% training and 20% testing data. The class balance in the training and testing data:

Training: 50.337% signal

Testing: 50.325% signal

D. Design and training of a 1D CNN classifier

1. Parameters in the 1D CNN:

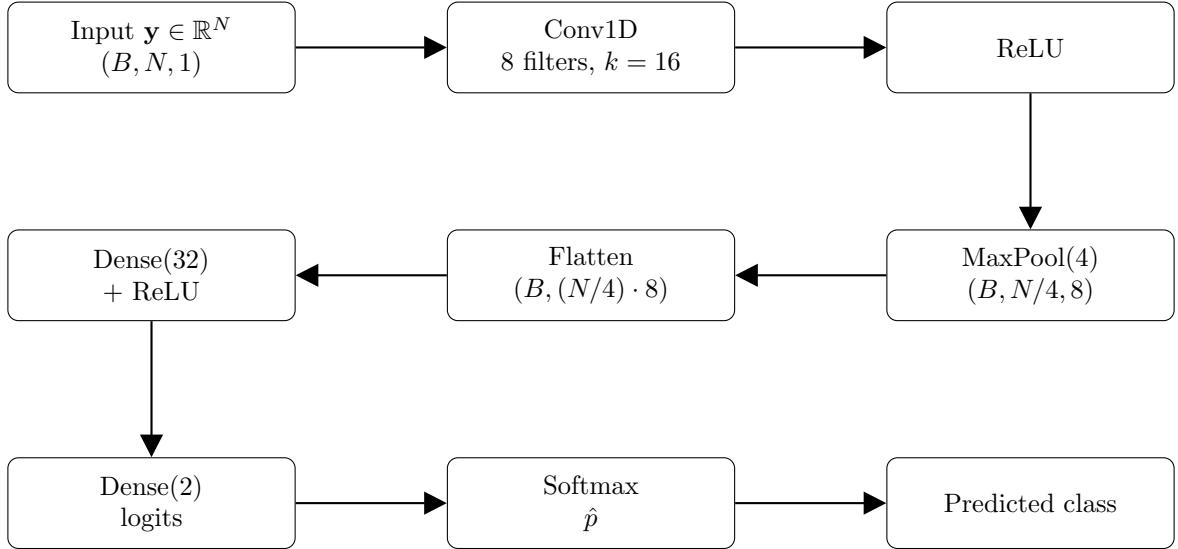


Figure 3: 1D CNN architecture for signal detection.

Conv1D layer ($1000 \rightarrow 1000$): $16 \times 1 \times 8 + 8 = 136$

Dense layer ($1000/4 \cdot 8 = 2000 \rightarrow 32$): $2000 \times 32 + 32 = 64032$

Dense output layer ($32 \rightarrow 2$): $32 \times 2 + 2 = 66$

Total trainable parameters = $136 + 64032 + 66 = 64234$

Total trainable parameters reported from listing 1 and 2 = 64234

2. The dataset is normalized by using the mean and standard deviation of training dataset.

$$\tilde{y}_i = \frac{y_i - \mu_{train}}{\sigma_{train}}$$

where, $i = 1, 2, \dots, N$.

Only training data statistics are used so that

- validation and test data do not influence the model by leaking information.
- optimization remains stabilized throughout the training.

Conv1D layer expects input of the shape (B, N, C) . So, $\mathbf{y} \in \mathbb{R}^N$ is reshaped to $(B, N, 1)$ (where, B is batch size) to add a channel dimension and make it compatible with convolution.

3. Please refer to listing 1 for the splitting and training of the data.

The set-up for training:

- Loss function: `softmax_cross_entropy_with_integer_labels()`
- Optimizer: `adamw()` (learning rate = 10^{-3} , weight decay = 5×10^{-4})
- Batch size: 256
- Maximum number of epochs: 20
- Maximum number of patience: 3

- Training stopped at 13th epoch as the early stopping criterion was met.

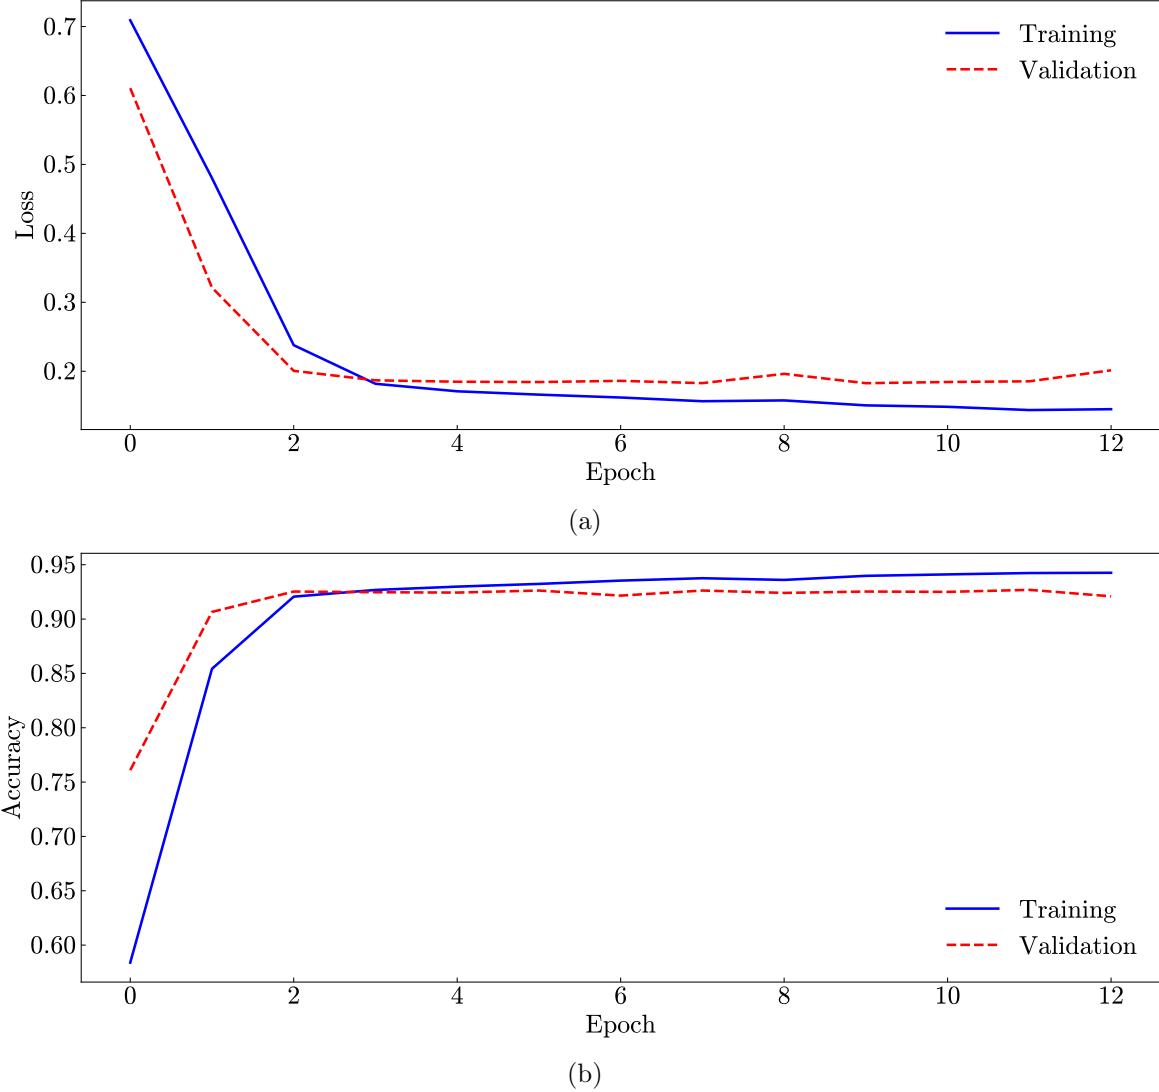


Figure 4: (a) Training and validation loss, and (b) training and validation accuracy of the model.

It is apparent from fig. 4 that the model started to overfit very mildly when the early stopping criterion (current validation loss $<$ best validation loss -10^{-6}) was activated, and training was stopped. The validation loss plateaus quickly within 3 epochs, and does not really change afterwards. The accuracy of the validation set also points to that. Overall, the model generalizes well with unprevailing tendency of overfitting.

- Test accuracy (threshold $\hat{\rho} = 0.5$): 0.9307500720024109 (final test loss: 0.1804021592993191)

E. Quantitative comparison: CNN vs matched filter

1. Please refer to listing 1.
2. AUC (matched filter) = 0.9840, AUC (CNN) = 0.9803.

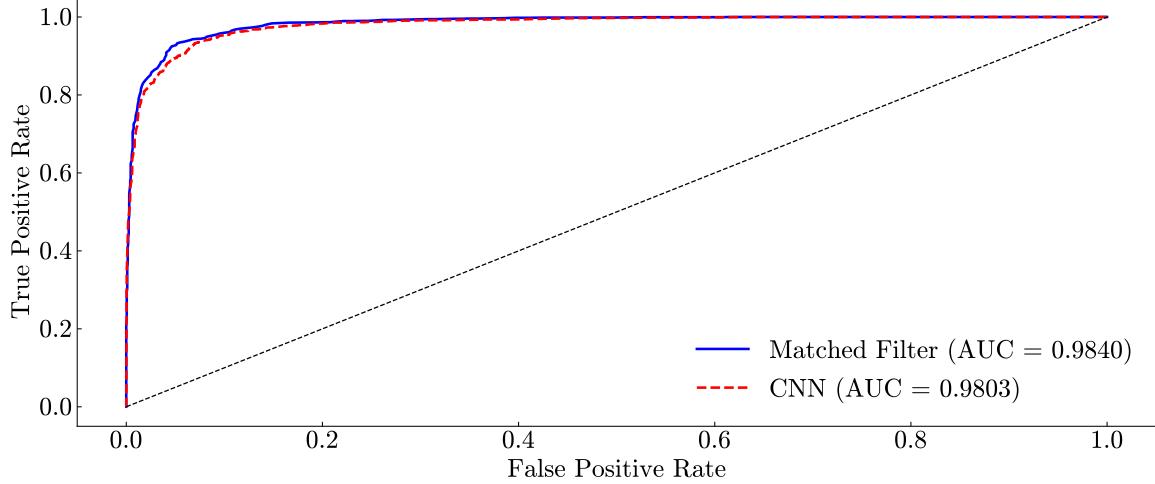


Figure 5: ROC curve for matched filter and CNN.

3. Discussions:

- CNN's ROC curve is very close to the ROC curve of matched filter. At very low FPR, matched filter has a higher peak compared to CNN, which is expected since matched filter is theoretically optimal. That is the only reason for AUC of CNN becoming 0.0037 less than matched filter's AUC. Overall, CNN impose almost the same level of performance as matched filter.
- At a few extremely low FPR values, CNN seems to have higher TPR than matched filter, but that is almost insignificant because it is more likely because of finite sampling noise rather than true outperformance. This is not surprising that matched filter proves to be superior as already mentioned that it is the optimal linear detector.

F. Going Further

• Interpreting the learned convolutional filters

1. From fig. 6, it is clear that none of the filters has any pattern that mimics template \hat{s} in terms of phase, period, and even amplitude. They look noisy and irregular, and does not look like smooth sinusoid. Filter 1 has the best correlation with the template, but the value of correlation is 0.3696 which is low enough to conclude that it did not rediscover a sinusoidal matched filter shape. The bottomline is that CNN seems to learn something different than matched filter despite having a performance that is very close to it.
2. CNN's learned filters are substantially different, since none of the kernels matches with template \hat{s} . So, CNN did not reproduce the optimal linear filter in Conv1. It seems

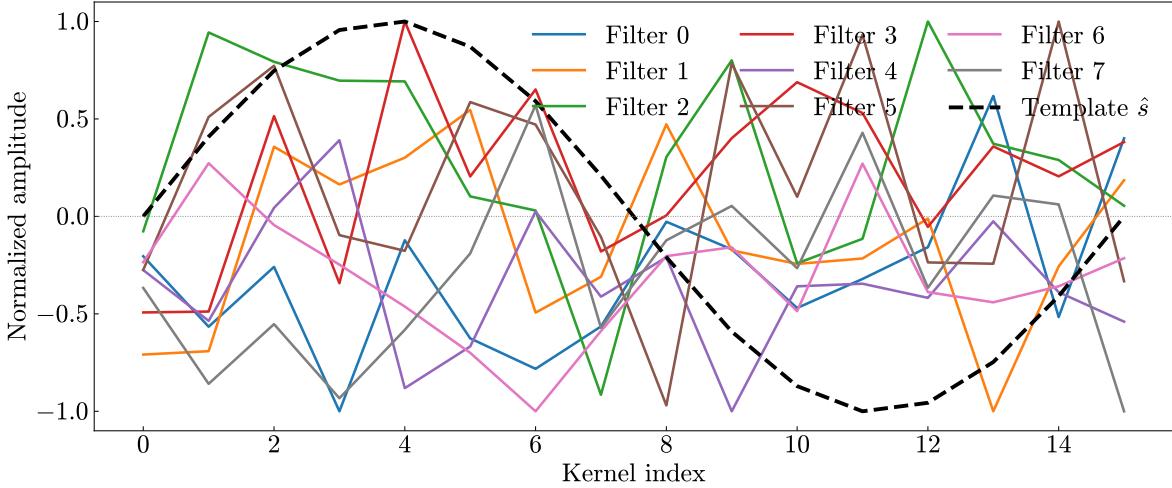
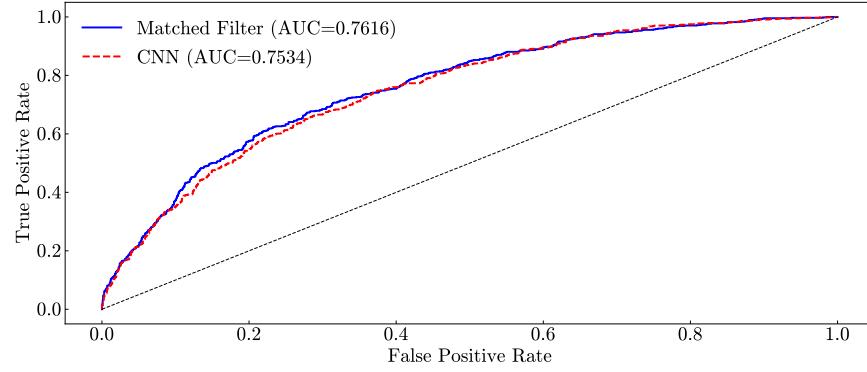


Figure 6: Learned 1D kernels and template \hat{s} as a function of time index.

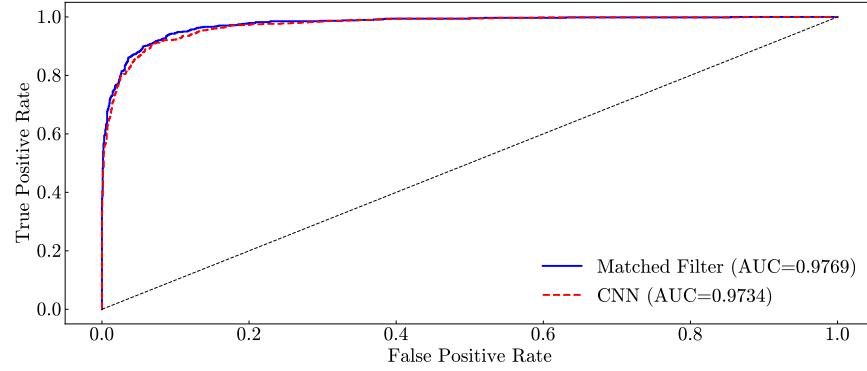
to pick on different features, likely constructing its classification rules through the non-linear combination of local features and dense layers.

- **Different signal strengths**

1. Please refer to listing 1.



(a)



(b)

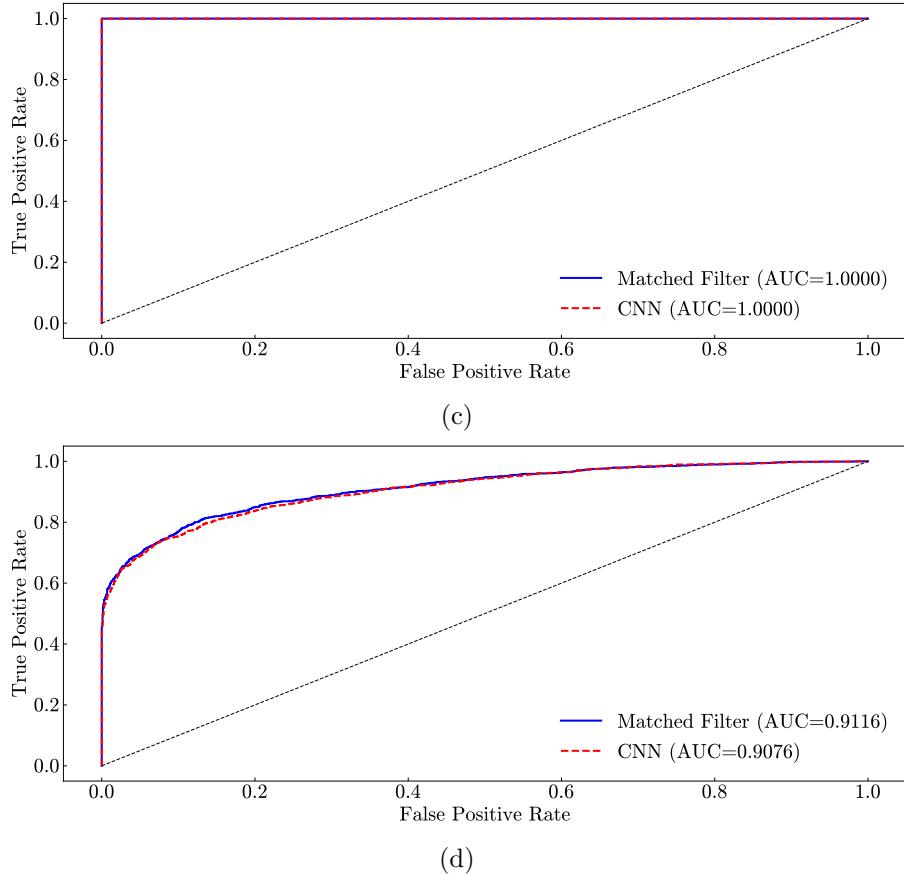


Figure 7: (a) ROC curves of matched filter and CNN for (a) $A = 1$, (b) $A = 3$, (c) $A = 10$, and (d) entire dataset.

2. From fig. 7, it is clear that matched filter remains superior as a class detector over CNN across all the amplitudes. However, the performance of CNN is very close to matched filter. One thing to note is that CNN can locally perform slightly better than matched filter as evident from fig. 7a and 7d clearly. Listing 2 shows a detailed statistics of the comparison between matched filter and CNN for all the amplitudes considered here.

```

1 import numpy as np
2 import jax
3 jax.config.update("jax_enable_x64", True)
4 import jax.numpy as jnp
5 import jax.random as random
6 import optax
7 from flax import linen as nn
8 from flax.core.frozen_dict import freeze, unfreeze
9 #from flax import nnx
10 from sklearn.model_selection import train_test_split
11 from sklearn.metrics import roc_curve, auc
12 import matplotlib.pyplot as plt
13 import matplotlib as mpl
14 from scipy.stats import norm
15
16 # parameters for plotting
17 plt.rcParams['font.family'] = 'serif'
18 plt.rcParams['font.serif'] = 'cmr10'
19 plt.rcParams['mathtext.fontset'] = 'cm'
20 plt.rcParams['font.size'] = 20
21 mpl.rcParams['axes.unicode_minus'] = False
22 plt.rcParams['axes.formatter.use_mathtext'] = True
23
24 def generate_data(A, N=1000, seed=None):
25     """Generate a noisy time series on {t_i, y_i}_{i=1}^N on [0, 2pi]"""
26     rng = np.random.default_rng(seed)
27     t = np.linspace(0.0, 2*np.pi, N)
28     dt = (2*np.pi) / (N - 1)
29     s_hat = np.sqrt(dt / np.pi) * np.sin(t)    # ||s_hat||_2 = 1
30     s = A * s_hat
31     n = rng.normal(loc=0.0, scale=1.0, size=N)
32     y = s + n
33     return t, y, s, n, s_hat, dt
34
35 def rho(y, s_hat):
36     """detection statistics"""
37     return np.dot(y, s_hat)
38
39 def sample_rhos(A, M, N=1000):
40     """draw M realizations of y at amplitude A and return rhos."""
41     rhos = np.zeros(M)
42     # s_hat is independent of A, so just grab it once
43     _, _, _, _, s_hat, _ = generate_data(A, N=N, seed=0)
44     for k in range(M):
45         _, y, _, _, _, _ = generate_data(A, N=N, seed=k)
46         rhos[k] = rho(y, s_hat)
47     return rhos, s_hat
48
49 # ===== B. Matched-filter recap and extended study =====
50 print("\n===== B. Matched-filter recap and extended study =====\n")
51 print("\n1. Generation of independent realizations of y:\n")
52 M = 20000
53 rhos_A0, s_hat = sample_rhos(A=0, M=M)
54 rhos_A3, _ = sample_rhos(A=3, M=M)
55
56 # plot for A=0
57 fig, ax = plt.subplots(figsize=(15, 6))
58 ax.hist(rhos_A0, bins=40, density=True, alpha=0.5, color='blue', edgecolor='black')

```

```

        , label='Empirical')
59 xs = np.linspace(-5, 5, 500)
60 ax.plot(xs, norm.pdf(xs, loc=0, scale=1), 'r-', label=rf'$\mathcal{N}(0, 1)$'
    )
61 plt.xlabel(r"$\rho$")
62 plt.ylabel("Density")
63 plt.legend(loc="upper right", frameon=False)
64 plt.tick_params(axis="both", which="both", direction="in")
65 plt.savefig("a0_pdf.pdf", dpi=1080)
66 plt.show()
67
68 # plot for A=3
69 fig, ax = plt.subplots(figsize=(15, 6))
70 ax.hist(rhos_A3, bins=40, density=True, alpha=0.5, color='blue', edgecolor='black',
    , label='Empirical')
71 xs = np.linspace(-2, 8, 500)
72 ax.plot(xs, norm.pdf(xs, loc=3, scale=1), 'r-', label=rf'$\mathcal{N}(3, 1)$'
    )
73 plt.xlabel(r"$\rho$")
74 plt.ylabel("Density")
75 plt.legend(loc="upper right", frameon=False)
76 plt.tick_params(axis="both", which="both", direction="in")
77 plt.savefig("a3_pdf.pdf", dpi=1080)
78 plt.show()
79
80 print("\n2. Creating filter_classifier and its verification:\n")
81 def filter_classifier(y, s_hat, rho0):
82     rho_val = np.dot(y, s_hat)
83     return int(rho_val >= rho0)
84
85 # get s_hat (does not depend on A)
86 t, _, _, _, s_hat, _ = generate_data(A=0, seed=0)
87
88 rho0 = 1.5 # threshold
89
90 # build dataset to predict class
91 def build_dataset(M=5000, N=1000, seed=0):
92     rng = np.random.default_rng(seed)
93     X = []
94     labels = []
95
96     for _ in range(M):
97         c = rng.integers(0, 2)      # 0 or 1
98         A = 0 if c == 0 else 3
99         _, y, _, _, _, _ = generate_data(A, N=N, seed=rng.integers(1e9))
100        X.append(y)
101        labels.append(c)
102
103    return np.array(X, dtype=np.float64), np.array(labels, dtype=np.int32)
104
105 X, labels = build_dataset(M=5000, N=1000, seed=0)
106
107 rhos = X @ s_hat
108 preds_theory = (rhos >= rho0).astype(int)
109
110 # compute classifier outputs
111 preds_func = np.array([filter_classifier(y, s_hat, rho0) for y in X])
112
113 print("All equal?", np.all(preds_theory == preds_func))

```

```

114 print("Mismatches:", np.sum(preds_theory != preds_func))
115
116 # plot for predicted class
117 fig, ax = plt.subplots(figsize=(15, 6))
118 ax.scatter(rhos, preds_func, s=30, color='blue', edgecolors='black', label='Classifier output')
119 ax.axvline(rho0, color='red', linestyle='--', linewidth=2, label=rf'$\rho_0 = \{rho0\}$')
120 plt.xlabel(r"$\rho$")
121 plt.ylabel("Predicted class")
122 plt.yticks([0,1])
123 plt.legend(frameon=False)
124 plt.tick_params(axis="both", which="both", direction="in")
125 plt.savefig("pred_class.pdf", dpi=1080)
126 plt.show()
127
128 # plot for PDF of predicted class
129 fig, ax = plt.subplots(figsize=(15, 6))
130 ax.hist(rhos[preds_func == 0], bins=40, density=True, alpha=0.5, edgecolor='black',
131         , color='blue', label='Predicted class 0')
132 ax.hist(rhos[preds_func == 1], bins=40, density=True, alpha=0.5, edgecolor='black',
133         , color='orange', label='Predicted class 1')
134 ax.axvline(rho0, color='red', linestyle='--', lw=2, label=rf'$\rho_0 = \{rho0\}$')
135 plt.xlabel(r"$\rho$")
136 plt.ylabel("Density")
137 plt.legend(frameon=False)
138 plt.tick_params(axis="both", which="both", direction="in")
139 plt.savefig("pred_class_pdf.pdf", dpi=1080)
140 plt.show()
141
142 # plot comparing theoretical prediction vs. function predictions
143 fig, ax = plt.subplots(figsize=(15, 6))
144 ax.scatter(np.arange(200), preds_theory[:200], facecolors='none', edgecolors='blue',
145         , s=60, linewidths=1.5, label='Theory')
146 ax.scatter(np.arange(200), preds_func[:200], c='red', marker='x', s=60, linewidths
147         =1.5, label='Function')
148 plt.xlabel("Sample index")
149 plt.yticks([0, 1])
150 plt.ylabel("Class")
151 plt.legend(frameon=False)
152 plt.tick_params(axis="both", which="both", direction="in")
153 plt.savefig("pred_class_comp.pdf", dpi=1080)
154 plt.show()
155
156 # ===== C. Building a labeled dataset for a CNN =====
157 print("\n===== C. Building a labeled dataset for a CNN =====\n")
158 def build_dataset(M, N, seed=0):
159     rng = np.random.default_rng(seed)
160     X = []
161     labels = []
162
163     for k in range(M):
164         c = rng.integers(0, 2)          # 0 = no signal, 1 = signal
165         A = 0 if c == 0 else 3
166         _, y, _, _, _, _ = generate_data(A, N=N, seed=rng.integers(1e9))
167         X.append(y)
168         labels.append(c)
169
170     X = np.array(X, dtype=np.float64)

```

```

167     labels = np.array(labels, dtype=np.int32)
168     return X, labels
169
170 print("\n1. Construction of a label dataset of time series:\n")
171 N = 1000
172 M = 20000
173 X, c = build_dataset(M=M, N=N, seed=0)
174
175 print(f"Class balance: {np.sum(c==0)} no-signal, {np.sum(c==1)} signal")
176 print(f"Percentage signal: {100*np.mean(c):.3f}%")
177
178 print("\n2. Split of dataset:\n")
179 # 80-20 split
180 X_train, X_test, c_train, c_test = train_test_split(
181     X, c,
182     test_size=0.20,
183     stratify=c,
184     random_state=0
185 )
186
187 print("Train:", X_train.shape, c_train.shape)
188 print("Test :", X_test.shape, c_test.shape)
189
190 print(f"Train class balance: {100*np.mean(c_train):.3f}% signal")
191 print(f"Test class balance: {100*np.mean(c_test):.3f}% signal")
192
193 # convert to JAX array
194 X_train = jnp.array(X_train, dtype=jnp.float64)
195 X_test = jnp.array(X_test, dtype=jnp.float64)
196 c_train = jnp.array(c_train)
197 c_test = jnp.array(c_test)
198
199 # ===== Design and training of a 1D CNN classifier =====
200 print("\n===== D. Design and training of a 1D CNN classifier =====\n")
201
202 # 1D CNN using flax linen
203 class CNN1D(nn.Module):
204     N: int
205
206     @nn.compact
207     def __call__(self, x):
208         # x is (B, N, 1)
209
210         # convolutional layer
211         x = nn.Conv(features=8, kernel_size=(16,), padding="SAME")(x)
212
213         # relu
214         x = nn.relu(x)
215
216         # max pooling
217         x = nn.max_pool(x, window_shape=(4,), strides=(4,))
218
219         # flatten
220         x = x.reshape((x.shape[0], -1)) # (B, 2000)
221
222         # one dense layer
223         x = nn.Dense(32)(x)
224         x = nn.relu(x)
225

```

```

226     # final output layer
227     logits = nn.Dense(2)(x)
228
229     return logits
230
231 print("\n1. 1D CNN architecture parameter counts:\n")
232 # functions for trainable parameter count
233 def count_params(params):
234     return sum(p.size for p in jax.tree_util.tree_leaves(params))
235
236 def detailed_param_count(params):
237     result = {}
238
239     def traverse(tree, parent=""):
240         for key, val in tree.items():
241             name = f"{parent}/{key}" if parent else key
242             if isinstance(val, dict):
243                 traverse(val, name)
244             else:
245                 result.setdefault(parent, 0)
246                 result[parent] += val.size
247
248     traverse(unfreeze(params))
249     return result
250
251 # initialize model and parameters
252 rng = jax.random.PRNGKey(0)
253 dummy = jnp.zeros((1, 1000, 1))    # NOTE: must be (B, N, 1)
254
255 model = CNN1D(N=1000)
256 params = model.init(rng, dummy)
257
258 # parameter counts
259 total = count_params(params)
260 print("Total parameters:", total)
261
262 details = detailed_param_count(params)
263 for layer, n in details.items():
264     print(f"{layer:20s} : {n}")
265
266 print("\n2. Normalize and reshape the data:\n")
267 # normalize the data
268 X_train_np = np.array(X_train)
269 X_test_np = np.array(X_test)
270
271 mean = X_train_np.mean()
272 std = X_train_np.std()
273
274 X_train_np = (X_train_np - mean) / std
275 X_test_np = (X_test_np - mean) / std
276
277 # reshape for Conv1D
278 X_train_np = X_train_np[:, :, None]
279 X_test_np = X_test_np[:, :, None]
280
281 print("\n3. Split for validation and training:\n")
282 # 80-20 train-validation split
283 X_tr_np, X_val_np, c_tr_np, c_val_np = train_test_split(
284     X_train_np, np.array(c_train), test_size=0.2,

```

```

285         stratify=np.array(c_train), random_state=1)
286
287 # convert to JAX arrays
288 X_tr = jnp.array(X_tr_np)
289 X_val = jnp.array(X_val_np)
290 X_te = jnp.array(X_test_np)
291
292 c_tr = jnp.array(c_tr_np)
293 c_val = jnp.array(c_val_np)
294 c_te = jnp.array(c_test)
295
296 print("\nTraining the 1D CNN:\n")
297
298 def loss_fn(params, batch_x, batch_y):
299     logits = model.apply(params, batch_x)
300     loss = optax.softmax_cross_entropy_with_integer_labels(logits, batch_y).mean()
301     return loss, logits
302
303 def accuracy(logits, labels):
304     preds = jnp.argmax(logits, axis=-1)
305     return jnp.mean(preds == labels)
306
307 optimizer = optax.adamw(learning_rate=1e-3, weight_decay=5e-4)
308
309 @jax.jit
310 def train_step(params, opt_state, xb, yb):
311     (loss, logits), grads = jax.value_and_grad(loss_fn, has_aux=True)(
312         params, xb, yb
313     )
314     updates, opt_state = optimizer.update(grads, opt_state, params)
315     params = optax.apply_updates(params, updates)
316     acc = accuracy(logits, yb)
317     return params, opt_state, loss, acc
318
319
320 @jax.jit
321 def eval_step(params, xb, yb):
322     loss, logits = loss_fn(params, xb, yb)
323     acc = accuracy(logits, yb)
324     return loss, acc
325
326 # initialize the model and optimizer
327 model = CNN1D(N=N)
328 rng = jax.random.PRNGKey(23)
329 params = model.init(rng, X_tr[:8])    # dummy batch for shapes
330 opt_state = optimizer.init(params)
331
332 # training params
333 batch_size = 256    # "None" for full batch training
334 epochs = 20
335
336 train_losses = []
337 train_accs = []
338 val_losses = []
339 val_accs = []
340
341 patience = 3
342 best_val_loss = jnp.inf
343 patience_counter = 0

```

```

344 | params_best = None
345 |
346 | # training loop
347 | for ep in range(epochs):
348 |
349 |     # reproducible shuffle
350 |     rng, sub = jax.random.split(rng)
351 |     perm = jax.random.permutation(sub, X_tr.shape[0])
352 |     X_shuf = X_tr[perm]
353 |     c_shuf = c_tr[perm]
354 |
355 |     # batching mode
356 |     if batch_size is None:
357 |         """full-batch"""
358 |         xb = X_shuf
359 |         yb = c_shuf
360 |
361 |         params, opt_state, loss, acc = train_step(params, opt_state, xb, yb)
362 |
363 |         ep_loss = float(loss)
364 |         ep_acc = float(acc)
365 |
366 |     else:
367 |         """mini-batch"""
368 |         num_batches = X_shuf.shape[0] // batch_size
369 |
370 |         ep_loss = 0.0
371 |         ep_acc = 0.0
372 |
373 |         for i in range(num_batches):
374 |             xb = X_shuf[i*batch_size:(i+1)*batch_size]
375 |             yb = c_shuf[i*batch_size:(i+1)*batch_size]
376 |
377 |             params, opt_state, loss, acc = train_step(
378 |                 params, opt_state, xb, yb
379 |             )
380 |             ep_loss += float(loss)
381 |             ep_acc += float(acc)
382 |
383 |         ep_loss /= num_batches
384 |         ep_acc /= num_batches
385 |
386 |     # validation
387 |     val_loss, val_acc = eval_step(params, X_val, c_val)
388 |
389 |     train_losses.append(ep_loss)
390 |     train_accs.append(ep_acc)
391 |     val_losses.append(float(val_loss))
392 |     val_accs.append(float(val_acc))
393 |
394 |     print(f"Epoch {ep+1:2d} | "
395 |           f"train_loss={ep_loss:.4f}, train_acc={ep_acc:.4f} | "
396 |           f"val_loss={float(val_loss):.4f}, val_acc={float(val_acc):.4f}")
397 |
398 |     # early stopping check
399 |     if float(val_loss) < float(best_val_loss) - 1e-6:
400 |         best_val_loss = float(val_loss)
401 |         params_best = params
402 |         patience_counter = 0

```

```

403     else:
404         patience_counter += 1
405
406     if patience_counter >= patience:
407         print(f"\nEarly stopping triggered at epoch {ep+1}")
408         params = params_best
409         break
410
411 # final test loss and accuracy
412 test_loss, test_acc = eval_step(params, X_te, c_te)
413 print("\nTest Loss:", float(test_loss))
414 print("Test Accuracy:", float(test_acc))
415
416 print("\n4. Plot training and validation loss:\n")
417 fig, ax = plt.subplots(figsize=(15, 6))
418 ax.plot(train_losses, "b-", lw=2, label="Training")
419 ax.plot(val_losses, "r--", lw=2, label="Validation")
420 plt.xlabel("Epoch")
421 plt.ylabel("Loss")
422 ax.legend(frameon=False)
423 ax.tick_params(axis="both", which="both", direction="in")
424 plt.savefig("cnn_loss.pdf", dpi=1080)
425 plt.show()
426
427 fig, ax = plt.subplots(figsize=(15, 6))
428 ax.plot(train_accs, "b-", lw=2, label="Training")
429 ax.plot(val_accs, "r--", lw=2, label="Validation")
430 plt.xlabel("Epoch")
431 plt.ylabel("Accuracy")
432 plt.legend(frameon=False)
433 plt.tick_params(axis="both", which="both", direction="in")
434 plt.savefig("cnn_acc.pdf", dpi=1080)
435 plt.show()
436
437 print("\n5. CNN's performance on the held out test set:\n")
438 test_logits = model.apply(params, X_te)
439 test_probs = jax.nn.softmax(test_logits, axis=-1)[:, 1]
440 test_preds = (test_probs >= 0.5).astype(int)
441
442 test_acc = jnp.mean(test_preds == c_te)
443
444 print("\nTest accuracy:", float(test_acc))
445
446 # ===== E. Quantitative comparison: CNN vs matched filter =====
447 print("\n===== E. Quantitative comparison: CNN vs matched filter =====\n")
448
449 print("\n1. Construct ROC:\n")
450
451 # matched-filter ROC
452 X_te_flat = np.array(X_te[..., 0])
453
454 rho_vals = X_te_flat @ s_hat
455
456 rho_range = np.linspace(rho_vals.min(), rho_vals.max(), 400)
457 tpr_mf = []
458 fpr_mf = []
459
460 for rho0 in rho_range:
461     preds = (rho_vals >= rho0).astype(int)

```

```

462
463     tp = np.sum((preds == 1) & (c_te == 1))
464     fp = np.sum((preds == 1) & (c_te == 0))
465     fn = np.sum((preds == 0) & (c_te == 1))
466     tn = np.sum((preds == 0) & (c_te == 0))
467
468     tpr_mf.append(tp / (tp + fn))
469     fpr_mf.append(fp / (fp + tn))
470
471 auc_mf = auc(fpr_mf, tpr_mf)
472 """
473 fpr_mf, tpr_mf, mf_thresholds = roc_curve(c_te, rho_vals)
474 auc_mf = auc(fpr_mf, tpr_mf)
475 """
476 # CNN ROC
477 """
478 def predict_cnn_proba(params, X):
479     logits = model.apply(params, X)
480     probs = jax.nn.softmax(logits, axis=-1)
481     return np.array(probs[:, 1]) # probability of class "1" (signal)
482
483 cnn_probs = predict_cnn_proba(params, X_te)
484
485 fpr_cnn, tpr_cnn, cnn_thresholds = roc_curve(c_te, cnn_probs)
486 auc_cnn = auc(fpr_cnn, tpr_cnn)
487 """
488 logits_te = model.apply(params, X_te)
489 probs = jax.nn.softmax(logits_te, axis=-1)
490 p_hat = np.array(probs[:, 1])
491
492 p_range = np.linspace(0, 1, 400)
493 tpr_cnn, fpr_cnn = [], []
494
495 for p0 in p_range:
496     preds = (p_hat >= p0).astype(int)
497
498     tp = np.sum((preds == 1) & (c_te == 1))
499     fp = np.sum((preds == 1) & (c_te == 0))
500     tn = np.sum((preds == 0) & (c_te == 0))
501     fn = np.sum((preds == 0) & (c_te == 1))
502
503     tpr_cnn.append(tp / (tp + fn))
504     fpr_cnn.append(fp / (fp + tn))
505
506 print("\n2. Plot ROC:\n")
507
508 fig, ax = plt.subplots(figsize=(15, 6))
509 ax.plot(fpr_mf, tpr_mf, 'b-', lw=2, label=f"Matched Filter (AUC = {auc_mf:.4f})")
510 ax.plot(fpr_cnn, tpr_cnn, 'r--', lw=2, label=f"CNN (AUC = {auc_cnn:.4f})")
511 ax.plot([0, 1], [0, 1], 'k--', lw=1)
512 plt.xlabel("False Positive Rate")
513 plt.ylabel("True Positive Rate")
514 plt.legend(frameon=False)
515 plt.tick_params(axis="both", which="both", direction="in")
516 plt.savefig("roc.pdf", dpi=1080)
517 plt.show()
518
519 # ===== F. Going further =====
520 print("\n===== F. Going further =====\n")

```

```

521 print("\nInterpreting the learned convolutional filters:\n")
522 p = params['params']
523
524 print("\nAvailable layers:", list(p.keys()))
525
526 # first convolutional layer
527 k1 = np.array(p['Conv_0']['kernel'])
528 print(f"Conv_0 kernel shape: {k1.shape}\n")
529 K1 = k1.shape[0] # kernel size
530 n_filters_1 = k1.shape[2] # number of learned filters
531
532 # extract each filter (input channel = 1)
533 kernels_conv1 = [k1[:, 0, j] for j in range(n_filters_1)]
534
535 # construct the known matched-filter template at same resolution
536 t_kernel = np.linspace(0, 2*np.pi, K1)
537 dt_kernel = 2*np.pi / (K1 - 1)
538 s_hat_kernel = np.sqrt(dt_kernel / np.pi) * np.sin(t_kernel)
539
540 # normalize for plotting
541 s_hat_norm = s_hat_kernel / np.max(np.abs(s_hat_kernel))
542
543 # plot for learned kernel and $|\hat{s}|$
544 fig, ax = plt.subplots(figsize=(15, 6))
545 for j, k_j in enumerate(kernels_conv1):
546     k_norm = k_j / (np.max(np.abs(k_j)) + 1e-12)
547     ax.plot(k_norm, lw=2, label=f"Filter {j}")
548 ax.plot(s_hat_norm, 'k--', lw=3, label="Template $|\hat{s}|$")
549 ax.axhline(0, color='gray', linestyle=':', linewidth=0.7)
550 plt.xlabel("Kernel index")
551 plt.ylabel("Normalized amplitude")
552 plt.legend(ncol=3, frameon=False)
553 plt.tick_params(axis="both", which="both", direction="in")
554 plt.savefig("comp_kernel.pdf", dpi=1080)
555 plt.show()
556
557 # correlation between learned filters and $|\hat{s}|$
558 s_norm = s_hat_kernel / np.linalg.norm(s_hat_kernel)
559
560 correlations = []
561 for j, k_j in enumerate(kernels_conv1):
562     k_norm = k_j / (np.linalg.norm(k_j) + 1e-12)
563     corr = np.dot(k_norm, s_norm)
564     correlations.append(corr)
565     print(f"Filter {j}: correlation = {corr:.4f}")
566
567 best_idx = np.argmax(np.abs(correlations))
568 best_corr = correlations[best_idx]
569
570 print(f"\nBest matching filter: Filter {best_idx}")
571 print(f"Absolute correlation: {abs(best_corr):.4f}")
572
573 print("\nDifferent signal strengths:")
574
575 # function for multi-amplitude dataset
576 def build_multiamp_dataset(M=20000, N=1000, seed=0):
577     rng = np.random.default_rng(seed)

```

```

580     X = []
581     labels = []
582     amplitudes = []
583
584     for k in range(M):
585         has_signal = rng.integers(0, 2) # 50-50 split
586
587         if has_signal == 0:
588             A = 0
589             label = 0 # no signal
590         else:
591             A = rng.choice([1, 3, 10]) # random signal strength
592             label = 1 # signal present (A >= 1)
593
594             _, y, _, _, _, _ = generate_data(A, N=N, seed=rng.integers(int(1e9)))
595             X.append(y)
596             labels.append(label)
597             amplitudes.append(A)
598
599     return np.array(X), np.array(labels), np.array(amplitudes)
600
601 X_multi, c_multi, A_multi = build_multiamp_dataset(M=20000, N=1000, seed=42)
602
603 # check distribution
604 print(f"\nDataset distribution:")
605 print(f"A=0: {np.sum(A_multi == 0)} ({100*np.mean(A_multi == 0):.1f}%)")
606 print(f"A=1: {np.sum(A_multi == 1)} ({100*np.mean(A_multi == 1):.1f}%)")
607 print(f"A=3: {np.sum(A_multi == 3)} ({100*np.mean(A_multi == 3):.1f}%)")
608 print(f"A=10: {np.sum(A_multi == 10)} ({100*np.mean(A_multi == 10):.1f}%)")
609
610 # 80-20 split
611 X_train_ma, X_test_ma, c_train_ma, c_test_ma, A_train_ma, A_test_ma =
612     train_test_split(
613         X_multi, c_multi, A_multi,
614         test_size=0.2,
615         stratify=c_multi,
616         random_state=0
617     )
618
619 # normalize test data
620 X_test_ma_norm = (X_test_ma - mean) / std
621 X_test_ma_jax = jnp.array(X_test_ma_norm, dtype=jnp.float64)[:, :, None]
622
623 # CNN predictions
624 cnn_logits_ma = model.apply(params, X_test_ma_jax)
625 cnn_probs_ma = np.array(jax.nn.softmax(cnn_logits_ma, axis=-1)[:, 1])
626
627 # matched filter predictions
628 rhos_test_ma = X_test_ma @ s_hat
629
630 # performance of CNN
631 fpr_cnn_all, tpr_cnn_all, _ = roc_curve(c_test_ma, cnn_probs_ma)
632 auc_cnn_all = auc(fpr_cnn_all, tpr_cnn_all)
633
634 # performance of matched filter
635 fpr_mf_all, tpr_mf_all, _ = roc_curve(c_test_ma, rhos_test_ma)
636 auc_mf_all = auc(fpr_mf_all, tpr_mf_all)
637
638 print(f"\nCNN AUC: {auc_cnn_all:.4f}")

```

```

638 print(f"Matched Filter AUC: {auc_mf_all:.4f}")
639
640 fig, ax = plt.subplots(figsize=(15, 6))
641 ax.plot(fpr_mf_all, tpr_mf_all, 'b-', lw=2, label=f'Matched Filter (AUC={auc_mf_all:.4f})')
642 ax.plot(fpr_cnn_all, tpr_cnn_all, 'r--', lw=2, label=f'CNN (AUC={auc_cnn_all:.4f})')
643 ax.plot([0, 1], [0, 1], 'k--', lw=1)
644 plt.xlabel('False Positive Rate')
645 plt.ylabel('True Positive Rate')
646 plt.legend(frameon=False)
647 plt.tick_params(axis="both", which="both", direction="in")
648 plt.savefig("comp_allamp.pdf", dpi=1080)
649 plt.show()
650
651 # per-amplitude analysis
652 for A_val in [1, 3, 10]:
653     # include both signal at this amplitude and all no-signal examples
654     mask = (A_test_ma == A_val) | (A_test_ma == 0)
655
656     print(f"\nFor A={A_val}:")
657     print(f"Total samples: {mask.sum()}")
658     print(f"A=0 (no signal): {np.sum(A_test_ma[mask] == 0)}")
659     print(f"A={A_val} (signal): {np.sum(A_test_ma[mask] == A_val)}")
660
661     # CNN ROC and AUC
662     fpr_cnn, tpr_cnn, _ = roc_curve(c_test_ma[mask], cnn_probs_ma[mask])
663     auc_cnn = auc(fpr_cnn, tpr_cnn)
664
665     # matched filter ROC and AUC
666     fpr_mf, tpr_mf, _ = roc_curve(c_test_ma[mask], rhos_test_ma[mask])
667     auc_mf = auc(fpr_mf, tpr_mf)
668
669     print(f"CNN AUC: {auc_cnn:.4f}")
670     print(f"Matched Filter AUC: {auc_mf:.4f}")
671     print(f"Difference: {abs(auc_cnn - auc_mf):.4f}")
672
673     fig, ax = plt.subplots(figsize=(15, 6))
674     ax.plot(fpr_mf, tpr_mf, 'b-', lw=2, label=f'Matched Filter (AUC={auc_mf:.4f})')
675     ax.plot(fpr_cnn, tpr_cnn, 'r--', lw=2, label=f'CNN (AUC={auc_cnn:.4f})')
676     ax.plot([0, 1], [0, 1], 'k--', lw=1)
677     plt.xlabel('False Positive Rate')
678     plt.ylabel('True Positive Rate')
679     plt.legend(frameon=False)
680     plt.tick_params(axis="both", which="both", direction="in")
681     plt.savefig(f"comp_amp{A_val}.pdf", dpi=1080)
682     plt.show()

```

Listing 1: cnn.py

```

1 ===== B. Matched-filter recap and extended study =====
2
3 1. Generation of independent realizations of y:
4
5 2. Creating filter_classifier and its verification:
6
7 All equal? True
8 Mismatches: 0
9
10 ===== C. Building a labeled dataset for a CNN =====
11
12 1. Construction of a label dataset of time series:
13
14 Class balance: 9933 no-signal, 10067 signal
15 Percentage signal: 50.335%
16
17 2. Split of dataset:
18
19 Train: (16000, 1000) (16000,)
20 Test : (4000, 1000) (4000,)
21 Train class balance: 50.337% signal
22 Test class balance: 50.325% signal
23
24 ===== D. Design and training of a 1D CNN classifier =====
25
26 1. 1D CNN architecture parameter counts:
27
28 Total parameters: 64234
29 params/Conv_0 : 136
30 params/Dense_0 : 64032
31 params/Dense_1 : 66
32
33 2. Normalize and reshape the data:
34
35 3. Split for validation and training:
36
37 Training the 1D CNN:
38
39 Epoch 1 | train_loss=0.7091, train_acc=0.5839 | val_loss=0.6106, val_acc=0.7609
40 Epoch 2 | train_loss=0.4802, train_acc=0.8542 | val_loss=0.3214, val_acc=0.9066
41 Epoch 3 | train_loss=0.2378, train_acc=0.9206 | val_loss=0.2006, val_acc=0.9253
42 Epoch 4 | train_loss=0.1818, train_acc=0.9268 | val_loss=0.1869, val_acc=0.9247
43 Epoch 5 | train_loss=0.1709, train_acc=0.9298 | val_loss=0.1848, val_acc=0.9244
44 Epoch 6 | train_loss=0.1661, train_acc=0.9323 | val_loss=0.1843, val_acc=0.9262
45 Epoch 7 | train_loss=0.1619, train_acc=0.9354 | val_loss=0.1861, val_acc=0.9216
46 Epoch 8 | train_loss=0.1565, train_acc=0.9376 | val_loss=0.1828, val_acc=0.9262
47 Epoch 9 | train_loss=0.1577, train_acc=0.9360 | val_loss=0.1962, val_acc=0.9241
48 Epoch 10 | train_loss=0.1505, train_acc=0.9398 | val_loss=0.1827, val_acc=0.9253
49 Epoch 11 | train_loss=0.1484, train_acc=0.9411 | val_loss=0.1844, val_acc=0.9250
50 Epoch 12 | train_loss=0.1437, train_acc=0.9423 | val_loss=0.1854, val_acc=0.9269
51 Epoch 13 | train_loss=0.1450, train_acc=0.9426 | val_loss=0.2014, val_acc=0.9209
52
53 Early stopping triggered at epoch 13
54
55 Test Loss: 0.18040215929931913
56 Test Accuracy: 0.9307500720024109
57
58 4. Plot training and validation loss:

```

```

59
60 5. CNN's performance on the held out test set:
61
62 Test accuracy: 0.9307500720024109
63
64 ===== E. Quantitative comparison: CNN vs matched filter =====
65
66 1. Construct ROC:
67
68 2. Plot ROC:
69
70 ===== F. Going further =====
71
72 Interpreting the learned convolutional filters:
73
74 Available layers: ['Conv_0', 'Dense_0', 'Dense_1']
75 Conv_0 kernel shape: (16, 1, 8)
76
77 Filter 0: correlation = -0.3457
78 Filter 1: correlation = +0.3696
79 Filter 2: correlation = +0.1077
80 Filter 3: correlation = -0.0548
81 Filter 4: correlation = +0.0973
82 Filter 5: correlation = +0.0004
83 Filter 6: correlation = -0.1656
84 Filter 7: correlation = -0.3674
85
86 Best matching filter: Filter 1
87 Absolute correlation: 0.3696
88
89 Different signal strengths:
90
91 Dataset distribution:
92 A=0: 9955 (49.8%)
93 A=1: 3399 (17.0%)
94 A=3: 3319 (16.6%)
95 A=10: 3327 (16.6%)
96
97 CNN AUC: 0.9076
98 Matched Filter AUC: 0.9116
99
100 For A=1:
101 Total samples: 2670
102 A=0 (no signal): 1991
103 A=1 (signal): 679
104 CNN AUC: 0.7534
105 Matched Filter AUC: 0.7616
106 Difference: 0.0082
107
108 For A=3:
109 Total samples: 2672
110 A=0 (no signal): 1991
111 A=3 (signal): 681
112 CNN AUC: 0.9734
113 Matched Filter AUC: 0.9769
114 Difference: 0.0035
115
116 For A=10:
117 Total samples: 2640

```

```
118 A=0 (no signal): 1991
119 A=10 (signal): 649
120 CNN AUC: 1.0000
121 Matched Filter AUC: 1.0000
122 Difference: 0.0000
```

Listing 2: Output terminal for `cnn.py`

II. NEURAL HAMILTONIAN FOR A 1D HARMONIC OSCILLATOR

B. Running the starter code & dataset generation

1.

$$k = 1, A = 1$$

- (a) $q_{num}(t)$ and $q_{exact}(t)$ are plotted in fig. 8 along with $p_{num}(t)$ and $p_{exact}(t)$.

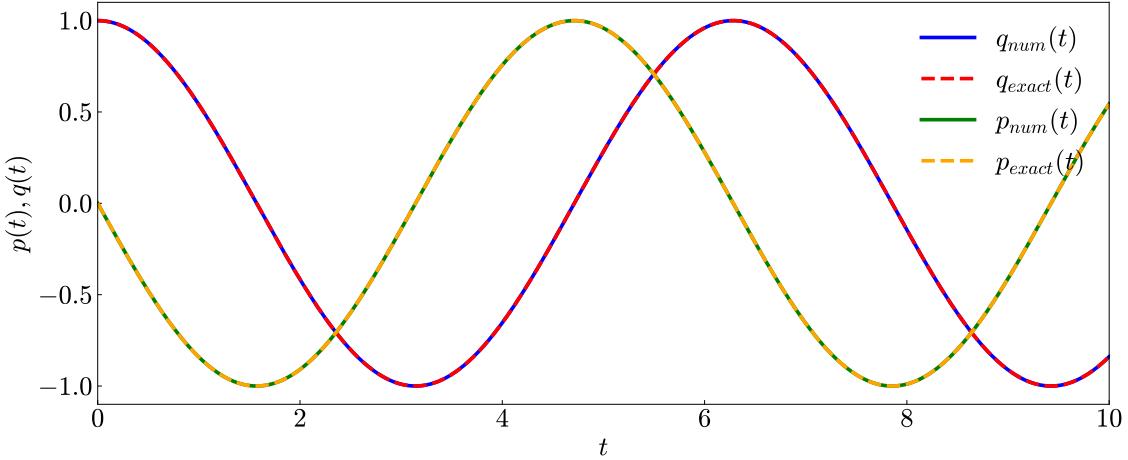


Figure 8: $q_{num}(t)$ and $q_{exact}(t)$.

They indeed look indistinguishable.

- (b) Error $e(t) = q_{num}(t) - q_{exact}(t)$ is plotted in fig. 9.

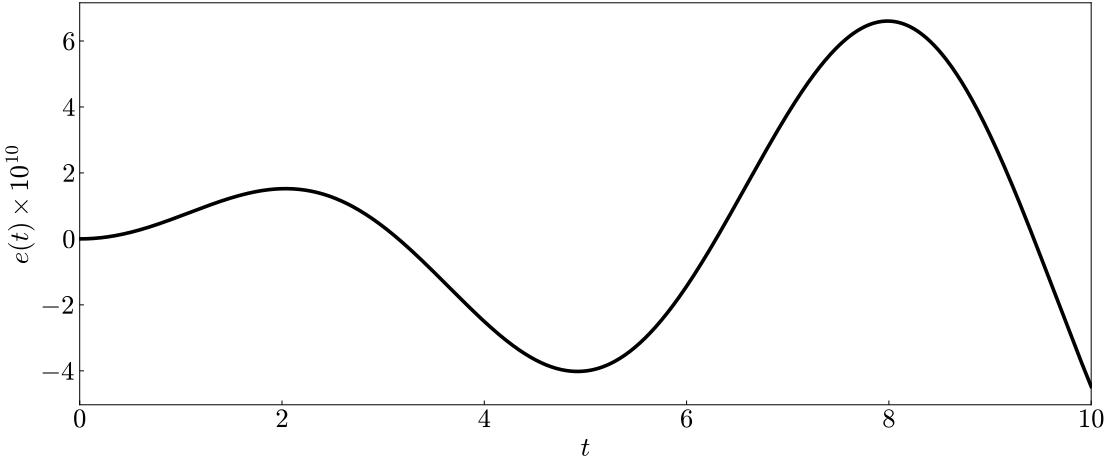


Figure 9: Error $e(t) = q_{num}(t) - q_{exact}(t)$.

The maximum absolute error as computed and reported in listing 3 and 4 respectively is 6.604×10^{-10} . The relative error with respect to $A = 1$ is also the same, which is very small. The difference is of course due to analytical and numerical way of solving the problem. However, the error is very tiny and insignificant.

- (c) So, the lower bound of error for the best possible learned model is 6.604×10^{-10} .

2.

$$k = 1, A \in \{0.5, 1.0, 1.5, 2.0\}$$

- (a) Please refer to listing 3.
- (b) Please refer to listing 3. (\dot{q}_n, \dot{p}_n) is obtained using analytic expressions:

$$\dot{q}_n = p_{exact}(t_n), \quad \dot{p}_n = -kq_{exact}(t_n)$$

- (c) Please refer to listing 3. Data points distribution in the 80-20 split:

Total: 4004, Training: 3203, Validation: 801

C. Learning dynamics with a neural Hamiltonian

- (a) Please refer to listing 3.

NN architecture:

- One hidden layer with 32 neurons.
- Activation function: `tanh`.

- (b) Please refer to listing 3.

- (c) Please refer to listing 3.

Hyperparameters:

- Optimizer: `lbfgs` (default settings).
- Maximum number of epochs: 500.
- Batch size: full batch.

- (d) Fig. 10 shows that the model does not overfit and validation loss keeps decreasing smoothly with decreasing training loss. So, the model generalizes well.

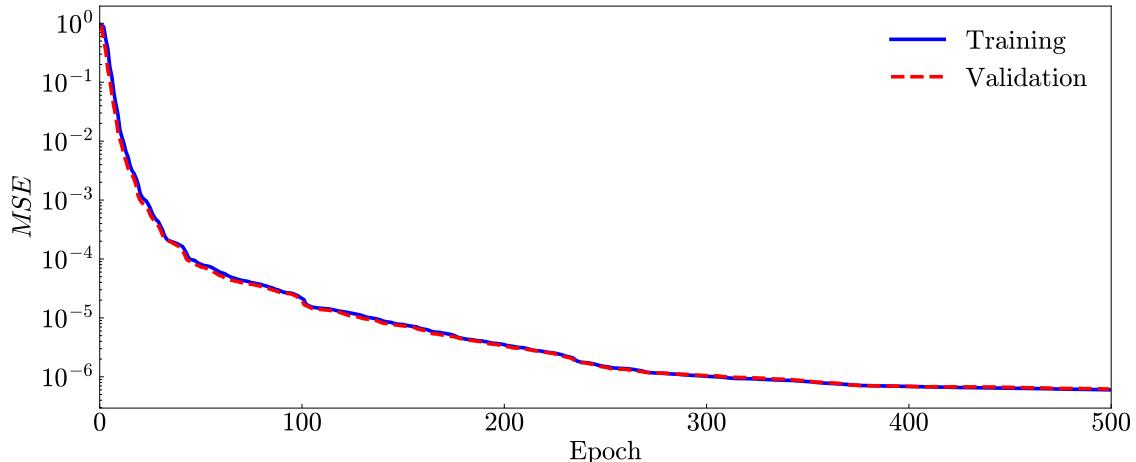


Figure 10: Training and validation losses vs. epoch.

D. Using your new model

- A_{test} is set at 1.25 for testing the model.

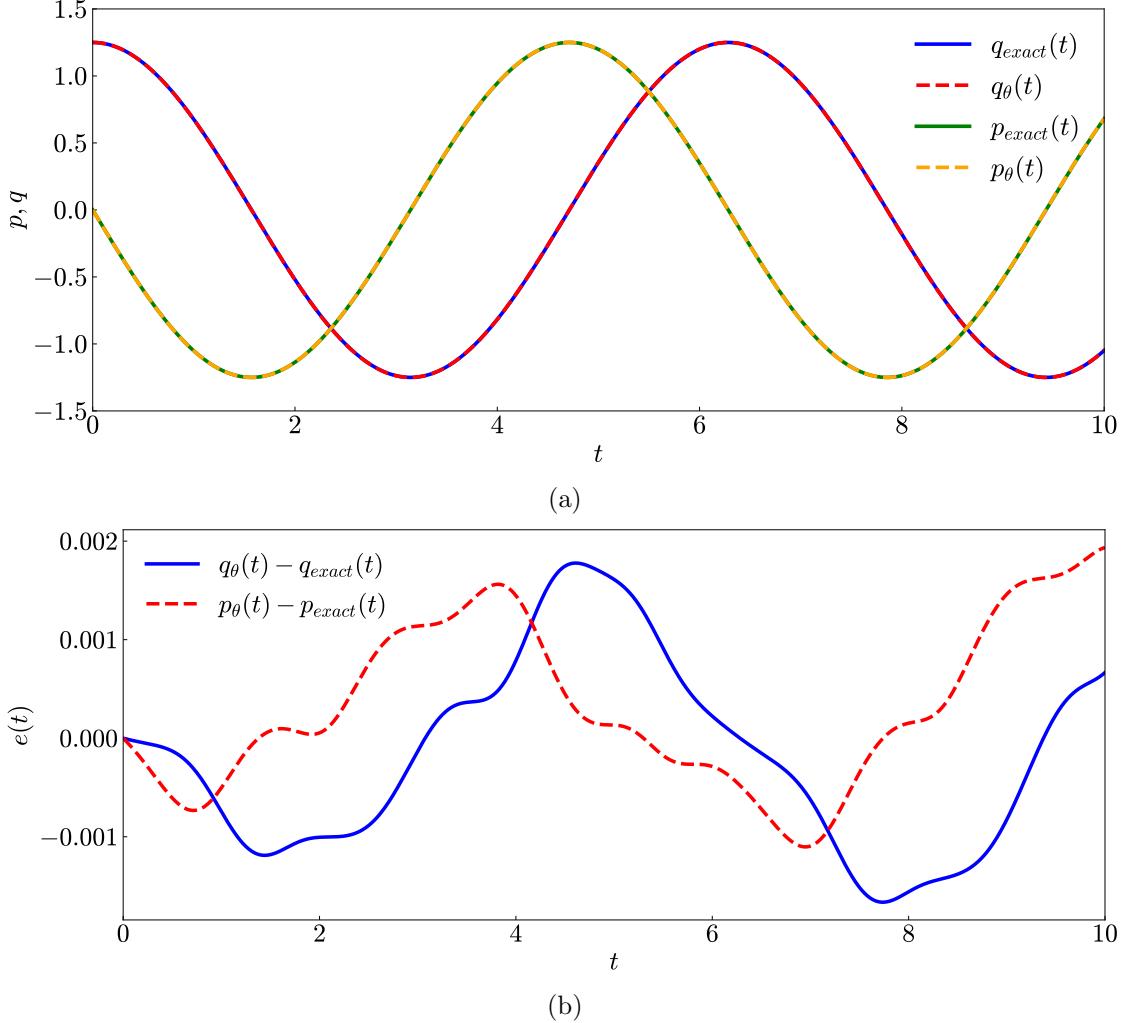


Figure 11: (a) Trajectories (q_θ, p_θ) and (q_{exact}, p_{exact}) , and (b) error between the trajectories.

Fig. 11 clearly shows that the model learns and performs really well for predicting the $H_\theta(q, p)$. The error bound between the model and exact solution is $\sim 10^{-3}$. So, the model generalizes really well, which was already evident from the training and validation losses in fig. 10.

- Clearly, the model predicts $H_\theta(q, p)$ well beyond $t = 10$ as can be seen from fig. 12. It has been run 10 times longer than it was trained for, still the error is $\sim 10^{-2}$ which further shows the fantastic generalization of the model.

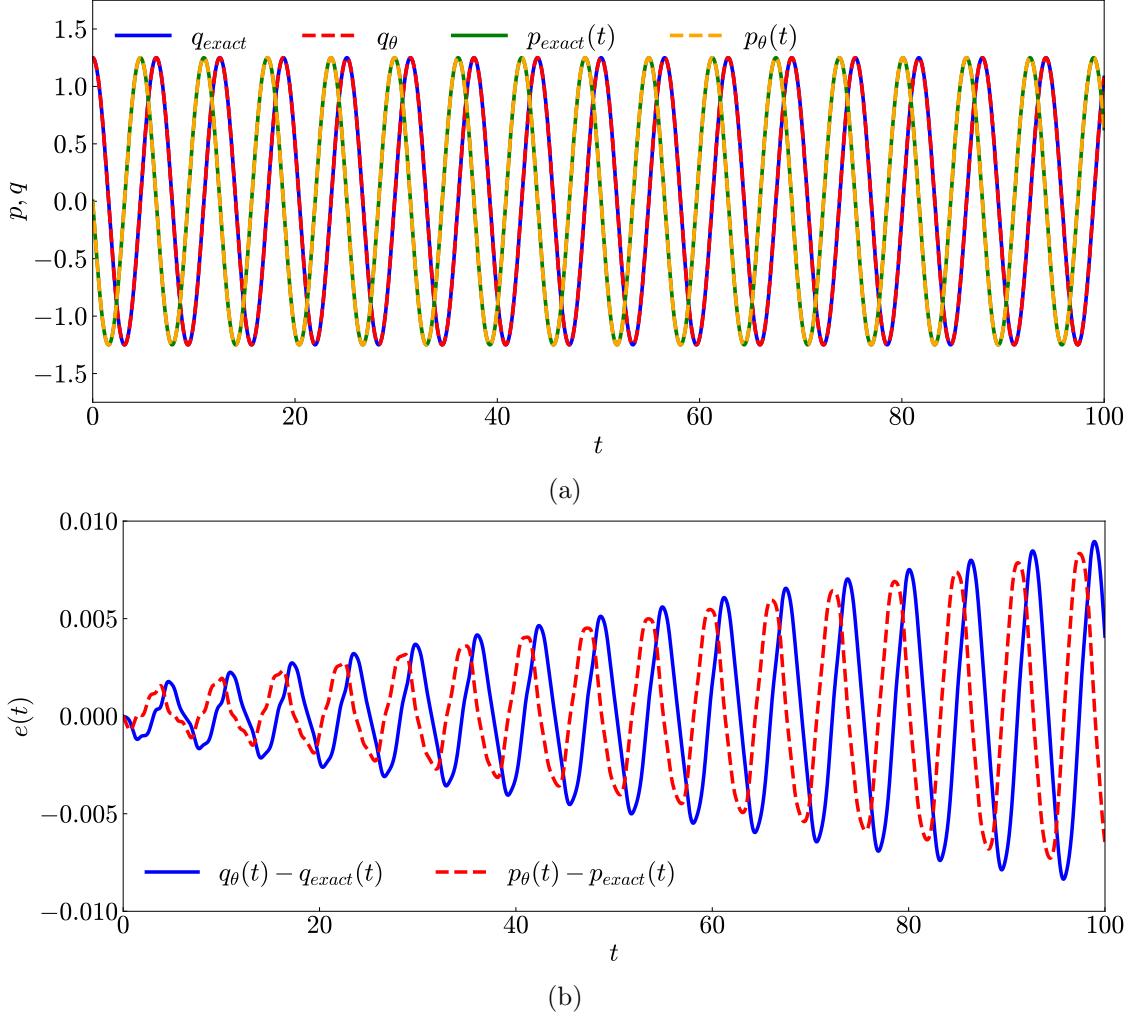


Figure 12: (a) Trajectories (q_θ, p_θ) and (q_{exact}, p_{exact}) , and (b) error between the trajectories extrapolated till $t = 100$.

E. Going Further

1. The Hamiltonian should be conserved over time, which means

$$\frac{dH}{dt} = 0$$

This can be checked using chain rule,

$$\frac{dH}{dt} = \frac{\partial H}{\partial q} \frac{dq}{dt} + \frac{\partial H}{\partial p} \frac{dp}{dt} = -\dot{p}\dot{q} + \dot{q}\dot{p} = 0$$

So, the conservation of energy, i.e., Hamiltonian can be enforced as a physics to the model.

2. The loss function then can be written as,

$$L_{\text{total},\theta} = \frac{1}{N_{\text{train}}} \sum_{n=1}^{N_{\text{train}}} \|\dot{q}_\theta(q_n, p_n) - \dot{q}_n\|^2 + \|\dot{p}_\theta(q_n, p_n) - \dot{p}_n\|^2 \\ + \lambda \frac{1}{N_{\text{train}} - 1} \sum_{n=1}^{N_{\text{train}}-1} \|H_\theta(q_{n+1}, p_{n+1}) - H_\theta(q_n, p_n)\|^2$$

where, $\lambda (= 10^{-5})$ is the penalization hyperparameter.

3. The model is re-run with the same hyperparameters as the data-only case.

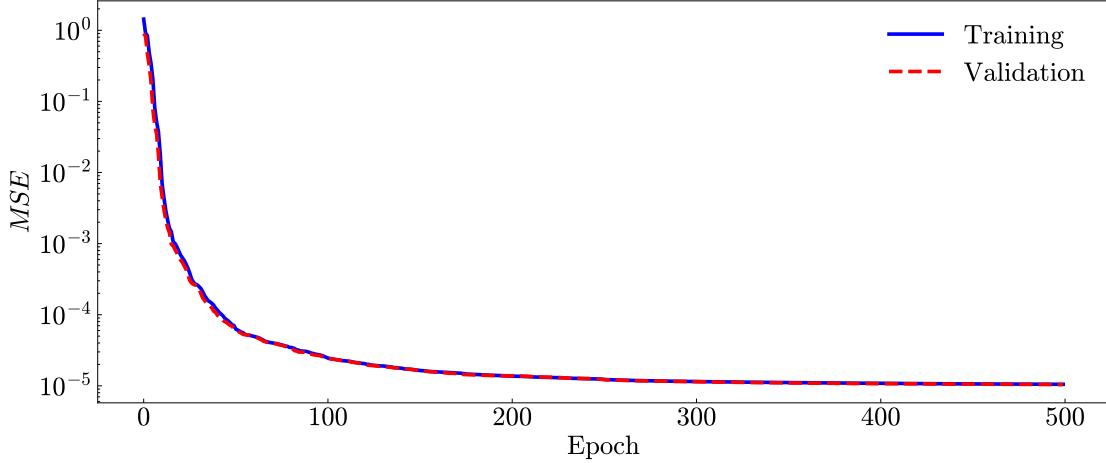
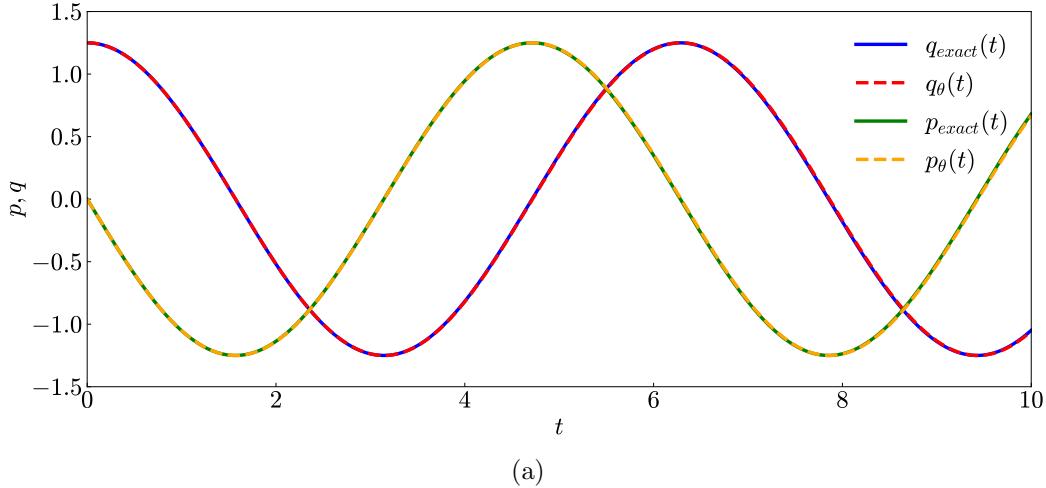


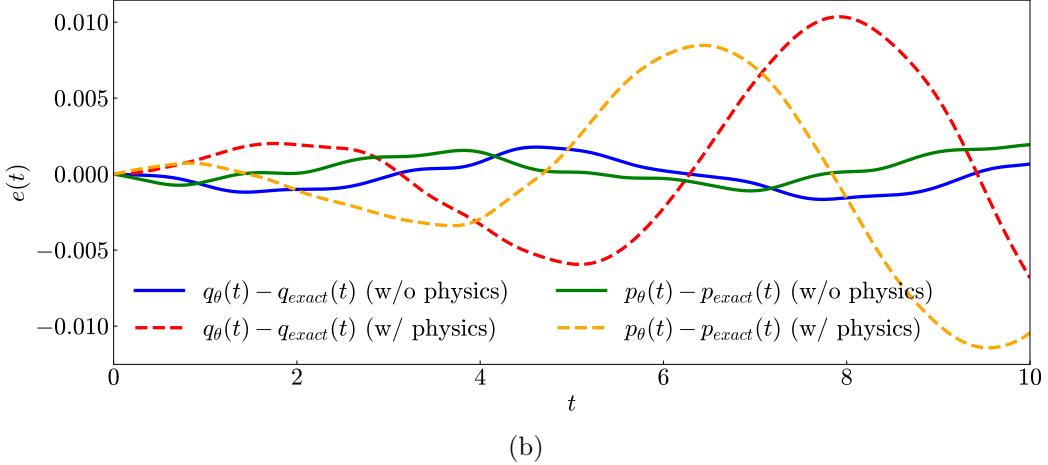
Figure 13: Training and validation losses w/ energy-penalty vs. epoch.

Fig. 13 shows that training and validation losses with physics incorporation are higher compared to fig. 10. Though the generalization is still good, but the loss is higher.

Fig. 14 and 15 show that the error in trajectories is higher with energy-penalty than without any penalty. The model can not perform as well as its data-only training when physics is

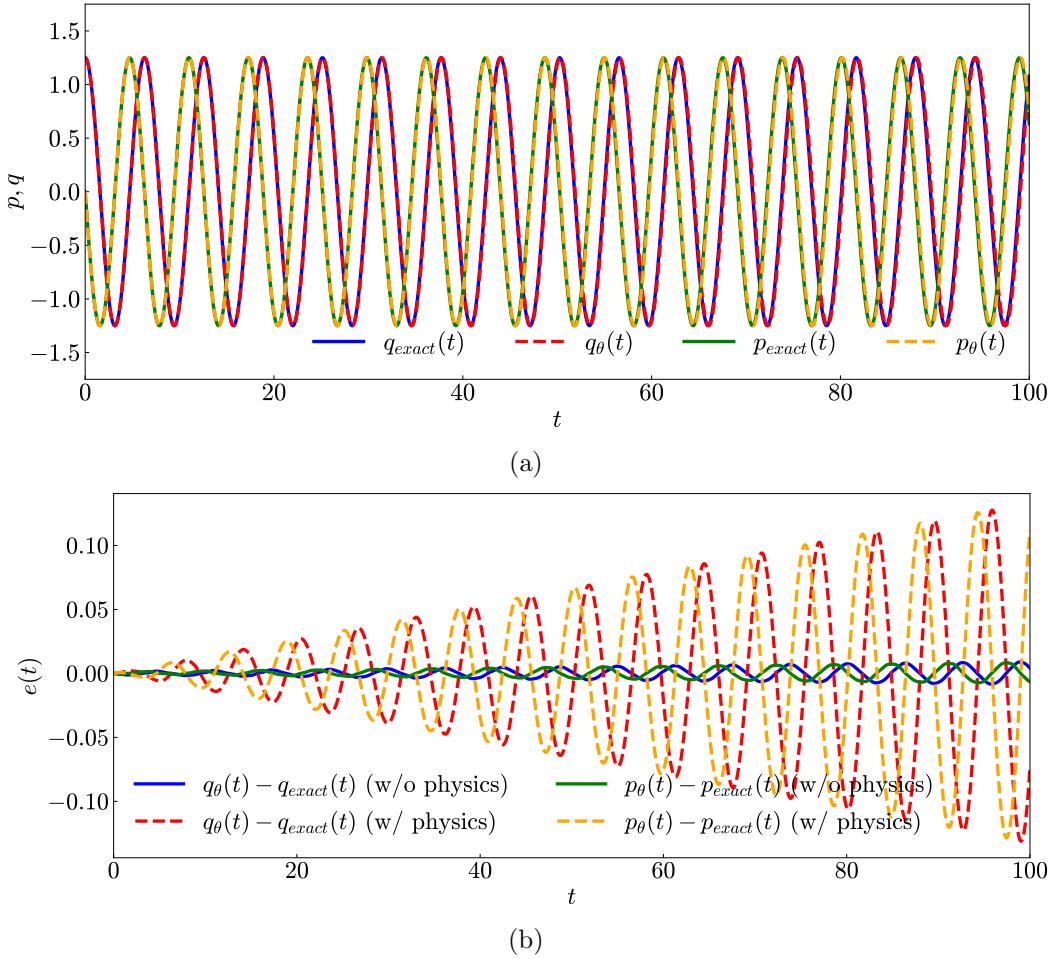


incorporated in the loss function. However, the overall generalization is still quite good.



(b)

Figure 14: (a) Trajectories (q_θ, p_θ) and (q_{exact}, p_{exact}) , and (b) error between the trajectories w/ energy-penalty.



(b)

Figure 15: (a) Trajectories (q_θ, p_θ) and (q_{exact}, p_{exact}) , and (b) error between the trajectories extrapolated till $t = 100$ w/ energy-penalty.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import matplotlib as mpl
4 import jax
5 jax.config.update("jax_enable_x64", True)
6 import optax
7 import jax.random as random
8 from flax import linen as nn
9 import jax.numpy as jnp
10 from typing import Sequence
11
12 # ===== B. Running the starter code & dataset generation =====
13 print("\n===== B. Running the starter code & dataset generation =====\n")
14
15 def sho_analytic(t, A=1.0, k=1.0):
16     """ Analytic solution for the 1D harmonic oscillator.
17
18     Initial conditions are
19     q(0) = A
20     p(0) = 0
21
22     Parameters are
23     mass m = 1
24     spring constant k > 0
25     """
26
27     omega = np.sqrt(k)
28     q = A * np.cos(omega*t)
29     p = -A*omega*np.sin(omega*t)
30     return q, p
31
32 def sho_rhs(q, p, k=1.0):
33     """Right-hand side of the
34
35     dq/dt = p
36     dp/dt = -k q
37     """
38
39     dqdt = p
40     dpdt = -k * q
41     return dqdt, dpdt
42
43 def rk4_step(q, p, dt, k=1.0):
44     """Single RK4 step taking us from time t to t+dt"""
45     k1_q, k1_p = sho_rhs(q, p, k)
46     k2_q, k2_p = sho_rhs(q + 0.5*dt*k1_q, p + 0.5*dt*k1_p, k)
47     k3_q, k3_p = sho_rhs(q + 0.5*dt*k2_q, p + 0.5*dt*k2_p, k)
48     k4_q, k4_p = sho_rhs(q + dt*k3_q, p + dt*k3_p, k)
49
50     q_next = q + (dt/6.0) * (k1_q + 2*k2_q + 2*k3_q + k4_q)
51     p_next = p + (dt/6.0) * (k1_p + 2*k2_p + 2*k3_p + k4_p)
52     return q_next, p_next
53
54 def sho_numerical(A=1.0, k=1.0, T=10.0, dt=0.01):
55     """ Numerical solution of the harmonic oscillator. """
56
57     t = np.arange(0.0, T + dt, dt)
58     q = np.zeros_like(t)

```

```

59     p = np.zeros_like(t)
60
61     # initial conditions
62     q[0] = A
63     p[0] = 0.0
64
65     for n in range(len(t) - 1):
66         q[n+1], p[n+1] = rk4_step(q[n], p[n], dt, k)
67
68     return t, q, p
69
70 print("\n1. Compute exact and numerical p, q:\n")
71
72 times, q_num, p_num = sho_numerical(A=1.0, k=1.0, T=10.0, dt=0.01)
73 q_exact, p_exact = sho_analytic(times, A=1.0, k=1.0)
74
75 # parameters for plotting
76 plt.rcParams['font.family'] = 'serif'
77 plt.rcParams['font.serif'] = 'cmr10'
78 plt.rcParams['mathtext.fontset'] = 'cm'
79 plt.rcParams['font.size'] = 22
80 mpl.rcParams['axes.unicode_minus'] = False
81 plt.rcParams['axes.formatter.use_mathtext'] = True
82
83 fig, ax = plt.subplots(figsize=(15, 6))
84 ax.plot(times, q_num, 'b-', lw=3, label=r'$q_{\text{num}}(t)$')
85 ax.plot(times, q_exact, 'r--', lw=3, label=r'$q_{\text{exact}}(t)$')
86 ax.plot(times, p_num, 'g-', lw=3, label=r'$p_{\text{num}}(t)$')
87 ax.plot(times, p_exact, '--', color="orange", lw=3, label=r'$p_{\text{exact}}(t)$')
88 plt.xlabel(r"$t$")
89 plt.ylabel(r"$p(t), q(t)$")
90 plt.xlim(0, 10)
91 plt.legend(frameon=False)
92 plt.tick_params(axis="both", which="both", direction="in")
93 plt.savefig("pq_exact_num.pdf", dpi=1080)
94 plt.show()
95
96 # plot error
97 fig, ax = plt.subplots(figsize=(15, 6))
98 ax.plot(times, (q_num - q_exact)*1e10, "k-", lw=3)
99 plt.xlabel(r"$t$")
100 plt.ylabel(r"$e(t) \times 10^{10}$")
101 plt.xlim(0, 10)
102 plt.tick_params(axis="both", which="both", direction="in")
103 plt.savefig("q_exact_num_err.pdf", dpi=1080)
104 plt.show()
105
106 max_err = np.max(np.abs(q_num - q_exact))
107 print(f"Maximum absolute error = {max_err:.3e}")
108
109 print("\n2. Data generation:\n")
110
111 k = 1.0
112 A_list = [0.5, 1.0, 1.5, 2.0]
113 T = 10.0
114 dt = 0.01
115
116 times = np.arange(0.0, T + dt, dt) # shape (1001,)
117 all_q = []

```

```

118 all_p = []
119
120 for A in A_list:
121     q_A, p_A = sho_analytic(times, A=A, k=k)
122     all_q.append(q_A)
123     all_p.append(p_A)
124
125 all_q = np.stack(all_q, axis=0) # shape (m, N_t)
126 all_p = np.stack(all_p, axis=0) # shape (m, N_t)
127
128 # flatten trajectories over all amplitudes and times
129 q_flat = all_q.reshape(-1) # shape (4004,)
130 p_flat = all_p.reshape(-1) # shape (4004,)
131
132 # use analytic expression (qdot = p_exact, pdot = -k * q_exact)
133 qdot_flat = p_flat.copy()
134 pdot_flat = -k * q_flat
135
136 # inputs X = (q, p), outputs y = (qdot, pdot)
137 X = np.stack([q_flat, p_flat], axis=1) # shape (4004, 2)
138 Y = np.stack([qdot_flat, pdot_flat], axis=1) # shape (4004, 2)
139
140 print("X shape:", X.shape)
141 print("Y shape:", Y.shape)
142
143 rng = np.random.default_rng(seed=0)
144
145 N = X.shape[0]
146 perm = rng.permutation(N)
147
148 # 80-20 split
149 train_frac = 0.8
150 N_train = int(train_frac * N)
151 train_idx = perm[:N_train]
152 val_idx = perm[N_train:]
153
154 X_train, Y_train = X[train_idx], Y[train_idx]
155 X_val, Y_val = X[val_idx], Y[val_idx]
156
157 print("Train size:", X_train.shape[0])
158 print("Val size : ", X_val.shape[0])
159
160 X_train_jax = jnp.asarray(X_train)
161 Y_train_jax = jnp.asarray(Y_train)
162 X_val_jax = jnp.asarray(X_val)
163 Y_val_jax = jnp.asarray(Y_val)
164
165 # ===== C. Learning dynamics with a neural Hamiltonian =====
166 print("\n===== C. Learning dynamics with a neural Hamiltonian =====\n")
167
168 # NN architecture
169 class HamiltonianNN(nn.Module):
170     hidden: Sequence[int] = (32,)
171
172     @nn.compact
173     def __call__(self, x):
174         x = x.astype(jnp.float64)
175         for h in self.hidden:
176             x = nn.tanh(nn.Dense(h, dtype=jnp.float64)(x))

```

```

177     #x = nn.relu(nn.Dense(h, dtype=jnp.float64)(x))
178     return nn.Dense(1, dtype=jnp.float64)(x)
179
180 # Hamiltonian vector fields' prediction
181 def hamiltonian_vector_field(params, model, qp):
182
183     def H_single(z):
184         return model.apply(params, z[None, :]).sum()
185
186     grads = jax.vmap(jax.grad(H_single))(qp)
187     dH_dq = grads[:, 0:1]
188     dH_dp = grads[:, 1:2]
189
190     return jnp.concatenate([dH_dp, -dH_dq], axis=1)
191
192 # loss function
193 def loss_fn(params, model, X, Y):
194     Y_pred = hamiltonian_vector_field(params, model, X)
195     return jnp.mean((Y_pred - Y)**2)
196
197 # lbfgs as optimizer
198 optimizer = optax.lbfgs()
199
200 model = HamiltonianNN()
201 key = random.PRNGKey(0)
202 params = model.init(key, X_train_jax[:1])
203 opt_state = optimizer.init(params)
204
205 def lbfgs_step(params, opt_state, X, Y):
206
207     # compute loss and gradient
208     loss, grad = jax.value_and_grad(loss_fn)(params, model, X, Y)
209
210     # update call
211     updates, opt_state = optimizer.update(
212         grad,
213         opt_state,
214         params=params,
215         value=loss,
216         grad=grad,
217         value_fn=lambda p: loss_fn(p, model, X, Y)
218     )
219
220     # apply update
221     params = optax.apply_updates(params, updates)
222     return params, opt_state, loss
223
224 # params
225 num_epochs = 500
226 train_losses = []
227 val_losses = []
228
229 # training loop
230 for epoch in range(1, num_epochs+1):
231
232     params, opt_state, train_loss = lbfgs_step(
233         params, opt_state, X_train_jax, Y_train_jax
234     )

```

```

236     val_loss = loss_fn(params, model, X_val_jax, Y_val_jax)
237
238     train_losses.append(train_loss)
239     val_losses.append(val_loss)
240
241     if epoch == 1 or epoch % 50 == 0:
242         print(f"Epoch {epoch}: train_loss={train_loss:.6e}, val_loss={val_loss:.6e}")
243
244 # plot training and validation losses
245 fig, ax = plt.subplots(figsize=(15, 6))
246 ax.semilogy(train_losses, "b-", lw=3, label="Training")
247 ax.semilogy(val_losses, "r--", lw=3, label="Validation")
248 plt.xlabel(r"Epoch")
249 plt.ylabel(r"$MSE$")
250 plt.xlim(0, 500)
251 plt.legend(frameon=False)
252 plt.tick_params(axis="both", which="both", direction="in")
253 plt.savefig("hamilton_loss.pdf", dpi=1080)
254 plt.show()
255
256 # ===== D. Using your new model =====
257 print("\n===== D. Using your new model =====\n")
258
259 A_test = 1.25
260
261 def learned_rhs(params, model, q, p):
262     qp = jnp.array([[q, p]])
263     qdot, pdot = hamiltonian_vector_field(params, model, qp)[0]
264     return float(qdot), float(pdot)
265
266 def rk4_step_learned(q, p, dt, params, model):
267     k1_q, k1_p = learned_rhs(params, model, q, p)
268     k2_q, k2_p = learned_rhs(params, model, q + 0.5*dt*k1_q, p + 0.5*dt*k1_p)
269     k3_q, k3_p = learned_rhs(params, model, q + 0.5*dt*k2_q, p + 0.5*dt*k2_p)
270     k4_q, k4_p = learned_rhs(params, model, q + dt*k3_q, p + dt*k3_p)
271     q_next = q + (dt/6.0)*(k1_q + 2*k2_q + 2*k3_q + k4_q)
272     p_next = p + (dt/6.0)*(k1_p + 2*k2_p + 2*k3_p + k4_p)
273     return q_next, p_next
274
275 def integrate_learned(A_test, T=10.0, dt=0.01):
276     t = np.arange(0.0, T+dt, dt)
277     q = np.zeros_like(t); p = np.zeros_like(t)
278     q[0] = A_test; p[0] = 0.0
279     for n in range(len(t)-1):
280         q[n+1], p[n+1] = rk4_step_learned(q[n], p[n], dt, params, model)
281     return t, q, p
282
283 # integrate and compare to analytic solution
284 t, q_learn, p_learn = integrate_learned(A_test, T=10.0, dt=0.01)
285 q_exact, p_exact = sho_analytic(t, A=A_test, k=1.0)
286
287 # plot trajectories
288 fig, ax = plt.subplots(figsize=(15, 6))
289 ax.plot(t, q_exact, 'b-', lw=3, label=r'$q_{exact}(t)$')
290 ax.plot(t, q_learn, 'r--', lw=3, label=r'$q_{\theta}(t)$')
291 ax.plot(t, p_exact, 'g-', lw=3, label=r'$p_{exact}(t)$')
292 ax.plot(t, p_learn, '--', color="orange", lw=3, label=r'$p_{\theta}(t)$')
293 plt.xlabel(r'$t$')

```

```

294 plt.ylabel(r'$p, q$')
295 plt.xlim(0,10)
296 plt.ylim(-1.5,1.5)
297 plt.legend(frameon=False)
298 plt.tick_params(axis="both", which="both", direction="in")
299 plt.savefig("learned_pq.pdf", dpi=1080)
300 plt.show()
301
302 fig, ax = plt.subplots(figsize=(15, 6))
303 ax.plot(t, q_learn - q_exact, 'b-', lw=3, label=r'$q_\theta(t) - q_{\text{exact}}(t)$')
304 ax.plot(t, p_learn - p_exact, 'r--', lw=3, label=r'$p_\theta(t) - p_{\text{exact}}(t)$')
305 plt.xlabel(r'$t$')
306 plt.ylabel(r'$e(t)$')
307 plt.xlim(0,10)
308 plt.legend(frameon=False)
309 plt.tick_params(axis="both", which="both", direction="in")
310 plt.savefig("learned_pq_err.pdf", dpi=1080)
311 plt.show()
312
313 # extrapolate
314 T_long = 100.0
315 t_long, q_learn_long, p_learn_long = integrate_learned(A_test, T=T_long, dt=0.01)
316 q_exact_long, p_exact_long = sho_analytic(t_long, A=A_test, k=1.0)
317
318 # plot trajectories over long period
319 fig, ax = plt.subplots(figsize=(15, 6))
320 ax.plot(t_long, q_exact_long, 'b-', lw=3, label=r'$q_{\text{exact}}$')
321 ax.plot(t_long, q_learn_long, 'r--', lw=3, label=r'$q_\theta$')
322 ax.plot(t_long, p_exact_long, 'g-', lw=3, label=r'$p_{\text{exact}}(t)$')
323 ax.plot(t_long, p_learn_long, '--', color="orange", lw=3, label=r'$p_\theta(t)$')
324 plt.xlabel(r'$t$')
325 plt.ylabel(r'$p, q$')
326 plt.xlim(t_long.min(), t_long.max())
327 plt.ylim(-1.75,1.75)
328 plt.legend(loc="upper left", ncol=4, frameon=False)
329 plt.tick_params(axis="both", which="both", direction="in")
330 plt.savefig("learned_pq_ex.pdf", dpi=1080)
331 plt.show()
332
333 # plot error vs time
334 fig, ax = plt.subplots(figsize=(15, 6))
335 ax.plot(t_long, q_learn_long - q_exact_long, "b-", lw=3, label=r'$q_\theta(t) - q_{\text{exact}}(t)$')
336 ax.plot(t_long, p_learn_long - p_exact_long, "r--", lw=3, label=r'$p_\theta(t) - p_{\text{exact}}(t)$')
337 plt.xlabel(r'$t$');
338 plt.ylabel(r'$e(t)$')
339 plt.xlim(t_long.min(), t_long.max())
340 plt.ylim(-0.01,0.01)
341 plt.legend(loc="lower left", ncol=2, frameon=False)
342 plt.tick_params(axis="both", which="both", direction="in")
343 plt.savefig("learned_pq_ex_err.pdf", dpi=1080)
344 plt.show()
345
346 # ===== E. Going further =====
347 # conservation-based physics loss
348 def loss_fn_phys(params, model, X_batch, Y_batch, lam=1e-5):
349     Y_pred = hamiltonian_vector_field(params, model, X_batch)
350     data_loss = jnp.mean((Y_pred - Y_batch)**2)

```

```

351
352     # energy conservation loss:  $H(q_{t+1}, p_{t+1}) - H(q_t, p_t)$  should be 0
353     X_t      = X_batch[:-1]
354     X_tp1   = X_batch[1:]
355
356     H_t      = model.apply(params, X_t)
357     H_tp1   = model.apply(params, X_tp1)
358
359     energy_conservation = jnp.mean((H_tp1 - H_t)**2)
360
361     return data_loss + lam * energy_conservation
362
363 def train_step_phys(params, opt_state, Xb, Yb):
364     loss, grads = jax.value_and_grad(loss_fn_phys)(params, model, Xb, Yb)
365
366     updates, opt_state = optimizer.update(
367         grads,
368         opt_state,
369         params=params,
370         value=loss,
371         grad=grads,
372         value_fn=lambda p: loss_fn_phys(p, model, Xb, Yb)
373     )
374     params = optax.apply_updates(params, updates)
375     return params, opt_state, loss
376
377
378 # retrain model
379 num_epochs = 500
380 train_losses = []
381 val_losses = []
382
383 params = model.init(random.PRNGKey(99), X_train_jax[:1])
384 opt_state = optimizer.init(params)
385
386 for epoch in range(1, num_epochs+1):
387     params, opt_state, train_loss = train_step_phys(
388         params, opt_state, X_train_jax, Y_train_jax
389     )
390     val_loss = loss_fn_phys(params, model, X_val_jax, Y_val_jax)
391
392     train_losses.append(train_loss)
393     val_losses.append(val_loss)
394
395     if epoch == 1 or epoch % 50 == 0:
396         print(f"Epoch {epoch}: train={train_loss:.3e}, val={val_loss:.3e}")
397
398 # plot physics-informed training performance
399 fig, ax = plt.subplots(figsize=(15, 6))
400 ax.semilogy(train_losses, "b-", lw=3, label="Training")
401 ax.semilogy(val_losses, "r--", lw=3, label="Validation")
402 plt.xlabel(r"Epoch")
403 plt.ylabel(r"$MSE$")
404 plt.legend(frameon=False)
405 plt.tick_params(axis="both", which="both", direction="in")
406 plt.savefig("hamilton_loss_phys.pdf", dpi=1080)
407 plt.show()
408
409 # evaluate for new amplitude

```

```

410 A_test = 1.25
411
412 def learned_rhs(params, model, q, p):
413     qp = jnp.array([[q, p]])
414     qdot, pdot = hamiltonian_vector_field(params, model, qp)[0]
415     return float(qdot), float(pdot)
416
417 def rk4_step_learned(q, p, dt):
418     k1_q, k1_p = learned_rhs(params, model, q, p)
419     k2_q, k2_p = learned_rhs(params, model, q + 0.5*dt*k1_q, p + 0.5*dt*k1_p)
420     k3_q, k3_p = learned_rhs(params, model, q + 0.5*dt*k2_q, p + 0.5*dt*k2_p)
421     k4_q, k4_p = learned_rhs(params, model, q + dt*k3_q, p + dt*k3_p)
422     return (
423         q + (dt/6)*(k1_q + 2*k2_q + 2*k3_q + k4_q),
424         p + (dt/6)*(k1_p + 2*k2_p + 2*k3_p + k4_p)
425     )
426
427 def integrate_learned(A, T, dt=0.01):
428     t = np.arange(0, T+dt, dt)
429     q = np.zeros_like(t); p = np.zeros_like(t)
430     q[0] = A; p[0] = 0
431     for n in range(len(t)-1):
432         q[n+1], p[n+1] = rk4_step_learned(q[n], p[n], dt)
433     return t, q, p
434
435 # for T = 10
436 t10, q_learn10, p_learn10 = integrate_learned(A_test, T=10)
437 q_exact10, p_exact10 = sho_analytic(t10, A=A_test)
438
439 # plot predicted trajectories
440 fig, ax = plt.subplots(figsize=(15, 6))
441 ax.plot(t10, q_exact10, 'b-', lw=3, label=r'$q_{\text{exact}}(t)$')
442 ax.plot(t10, q_learn10, 'r--', lw=3, label=r'$q_{\text{learn}}(t)$')
443 ax.plot(t10, p_exact10, 'g--', lw=3, label=r'$p_{\text{exact}}(t)$')
444 ax.plot(t10, p_learn10, '--', color="orange", lw=3, label=r'$p_{\text{learn}}(t)$')
445 plt.xlabel(r'$t$')
446 plt.ylabel(r'$p, q$')
447 plt.xlim(t10.min(), t10.max())
448 plt.ylim(-1.5, 1.5)
449 plt.legend(frameon=False)
450 plt.tick_params(axis="both", which="both", direction="in")
451 plt.savefig("learned_pq_phys.pdf", dpi=1080)
452 plt.show()
453
454 fig, ax = plt.subplots(figsize=(15, 6))
455 ax.plot(t, q_learn - q_exact, 'b-', lw=3, label=r'$q_{\text{learn}} - q_{\text{exact}}(t)$ (w/o physics)')
456 ax.plot(t, q_learn10 - q_exact10, 'r--', lw=3, label=r'$q_{\text{learn}} - q_{\text{exact}}(t)$ (w/ physics)')
457 ax.plot(t, p_learn - p_exact, 'g--', lw=3, label=r'$p_{\text{learn}} - p_{\text{exact}}(t)$ (w/o physics)')
458 ax.plot(t, p_learn10 - p_exact10, '--', color="orange", lw=3, label=r'$p_{\text{learn}} - p_{\text{exact}}(t)$ (w/ physics)')
459 plt.xlabel(r'$t$')
460 plt.ylabel(r'$e(t)$')
461 plt.xlim(t10.min(), t10.max())
462 plt.legend(ncol=2, frameon=False)
463 plt.tick_params(axis="both", which="both", direction="in")
464 plt.savefig("learned_pq_err_phys.pdf", dpi=1080)

```

```

465 plt.show()
466
467 # extrapolation
468 t_long100, q_learn_long100, p_learn_long100 = integrate_learned(A_test, T=100)
469 q_exact_long100, p_exact_long100 = sho_analytic(t_long100, A=A_test)
470
471 fig, ax = plt.subplots(figsize=(15, 6))
472 ax.plot(t_long100, q_exact_long100, 'b-', lw=3, label=r'$q_{\text{exact}}(t)$')
473 ax.plot(t_long100, q_learn_long100, 'r--', lw=3, label=r'$q_{\text{learn}}(t)$')
474 ax.plot(t_long100, p_exact_long100, 'g-', lw=3, label=r'$p_{\text{exact}}(t)$')
475 ax.plot(t_long100, p_learn_long100, '--', lw=3, color="orange", label=r'$p_{\text{learn}}(t)$')
476 plt.xlabel(r'$t$');
477 plt.ylabel(r'$p, q$')
478 plt.xlim(t_long100.min(), t_long100.max())
479 plt.ylim(-1.75, 1.75)
480 plt.legend(ncol=4, frameon=False)
481 plt.tick_params(axis="both", which="both", direction="in")
482 plt.savefig("learned_pq_ex_phys.pdf", dpi=1080)
483 plt.show()
484
485 fig, ax = plt.subplots(figsize=(15, 6))
486 ax.plot(t_long, q_learn_long - q_exact_long, 'b-', lw=3, label=r'$q_{\text{learn}} - q_{\text{exact}}(t)$ (w/o physics)')
487 ax.plot(t_long100, q_learn_long100 - q_exact_long100, 'r--', lw=3, label=r'$q_{\text{learn}} - q_{\text{exact}}(t)$ (w/ physics)')
488 ax.plot(t_long, p_learn_long - p_exact_long, 'g-', lw=3, label=r'$p_{\text{learn}} - p_{\text{exact}}(t)$ (w/o physics)')
489 ax.plot(t_long100, p_learn_long100 - p_exact_long100, '--', lw=3, color="orange", label=r'$p_{\text{learn}} - p_{\text{exact}}(t)$ (w/ physics)')
490 plt.xlabel(r'$t$');
491 plt.ylabel(r'$e(t)$')
492 plt.xlim(t_long100.min(), t_long100.max())
493 plt.legend(ncol=2, frameon=False)
494 plt.tick_params(axis="both", which="both", direction="in")
495 plt.savefig("learned_pq_ex_err_phys.pdf", dpi=1080)
496 plt.show()

```

Listing 3: neural_hamiltonian.py

```

1 ===== B. Running the starter code & dataset generation =====
2
3 1. Compute exact and numerical p, q:
4
5 Maximum absolute error = 6.604e-10
6
7 2. Data generation:
8
9 X shape: (4004, 2)
10 Y shape: (4004, 2)
11 Train size: 3203
12 Val size : 801
13
14 ===== C. Learning dynamics with a neural Hamiltonian =====
15
16 Epoch 1: train_loss=9.587591e-01, val_loss=9.098519e-01
17 Epoch 50: train_loss=8.340965e-05, val_loss=7.641729e-05
18 Epoch 100: train_loss=2.237719e-05, val_loss=2.110202e-05
19 Epoch 150: train_loss=7.755750e-06, val_loss=7.434588e-06
20 Epoch 200: train_loss=3.566024e-06, val_loss=3.408729e-06
21 Epoch 250: train_loss=1.507338e-06, val_loss=1.473887e-06
22 Epoch 300: train_loss=1.030428e-06, val_loss=1.050491e-06
23 Epoch 350: train_loss=8.345358e-07, val_loss=8.522308e-07
24 Epoch 400: train_loss=6.878314e-07, val_loss=6.896047e-07
25 Epoch 450: train_loss=6.419417e-07, val_loss=6.636466e-07
26 Epoch 500: train_loss=6.012402e-07, val_loss=6.222261e-07
27
28 ===== D. Using your new model =====
29
30 ===== E. Going further =====
31
32 Epoch 1: train=1.435e+00, val=9.000e-01
33 Epoch 50: train=7.056e-05, val=6.243e-05
34 Epoch 100: train=2.509e-05, val=2.466e-05
35 Epoch 150: train=1.658e-05, val=1.643e-05
36 Epoch 200: train=1.371e-05, val=1.384e-05
37 Epoch 250: train=1.237e-05, val=1.231e-05
38 Epoch 300: train=1.151e-05, val=1.146e-05
39 Epoch 350: train=1.111e-05, val=1.099e-05
40 Epoch 400: train=1.086e-05, val=1.074e-05
41 Epoch 450: train=1.063e-05, val=1.054e-05
42 Epoch 500: train=1.051e-05, val=1.042e-05

```

Listing 4: Output terminal for `neural_hamiltonian.py`